

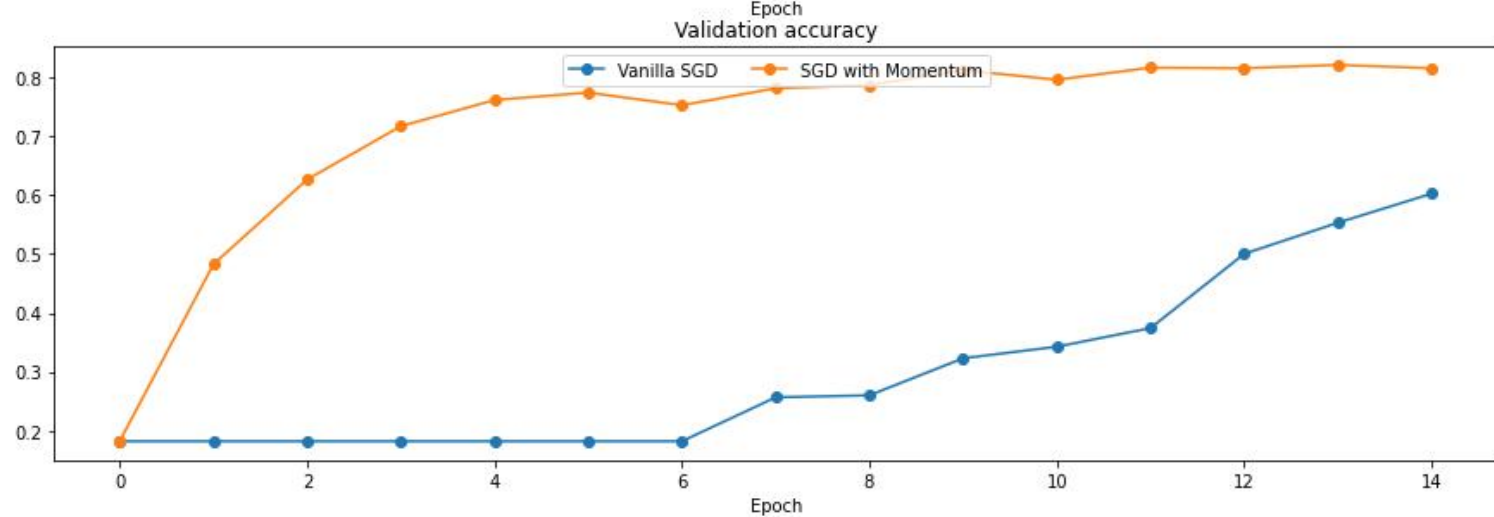
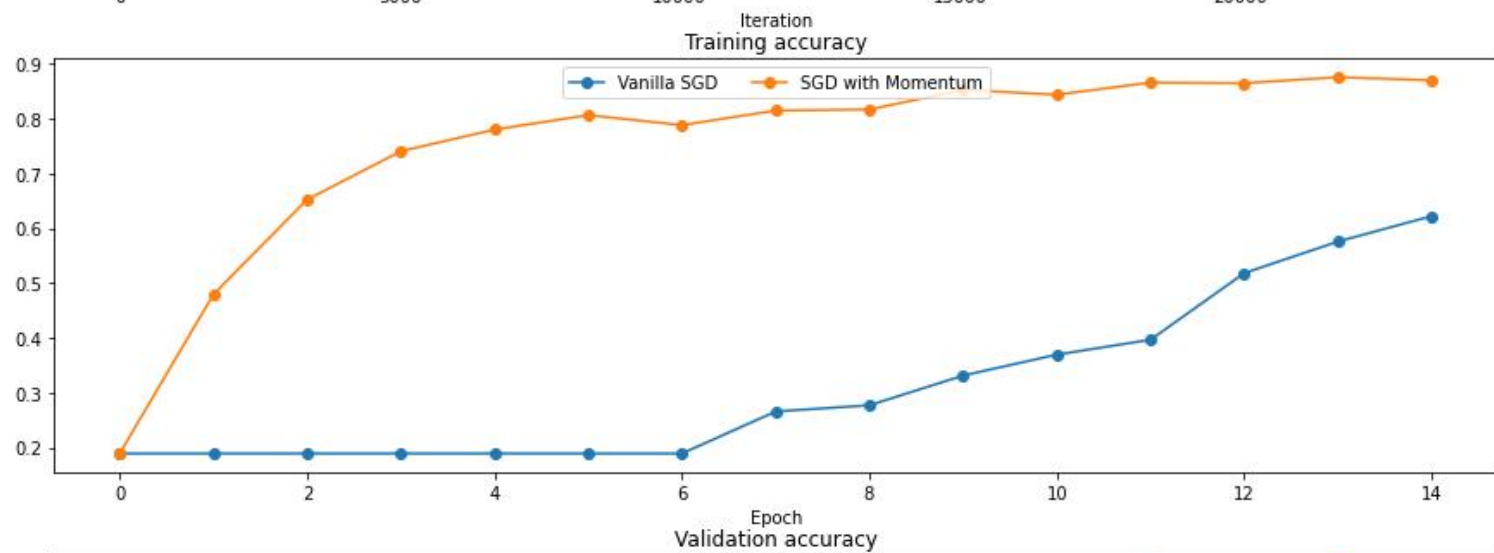


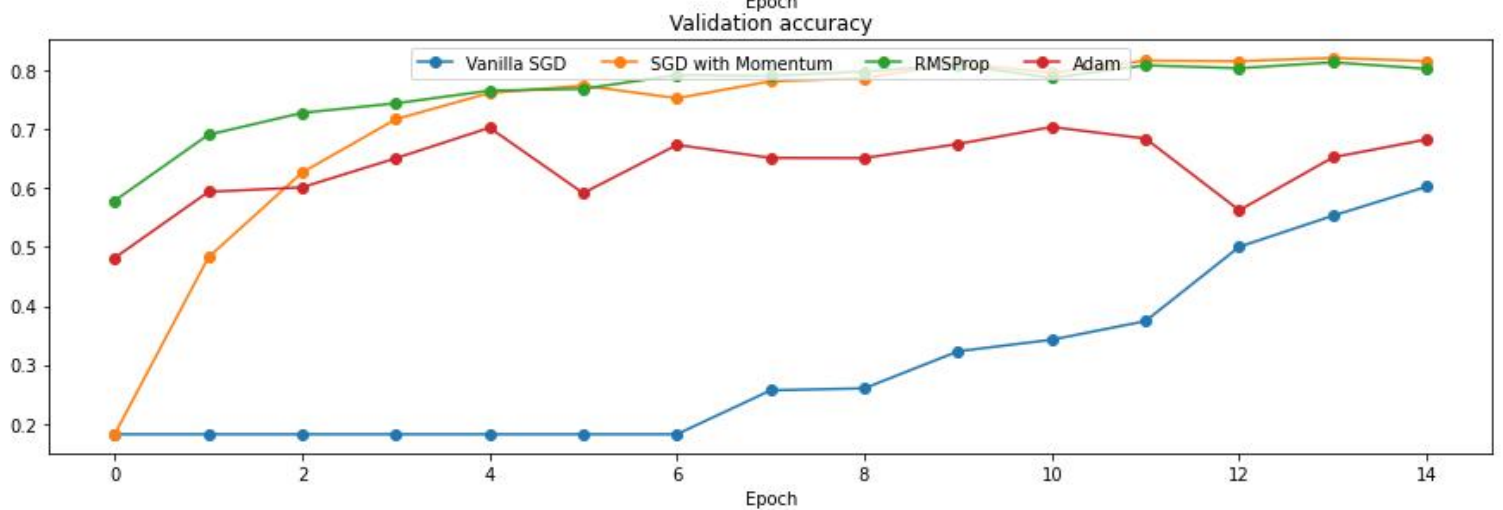
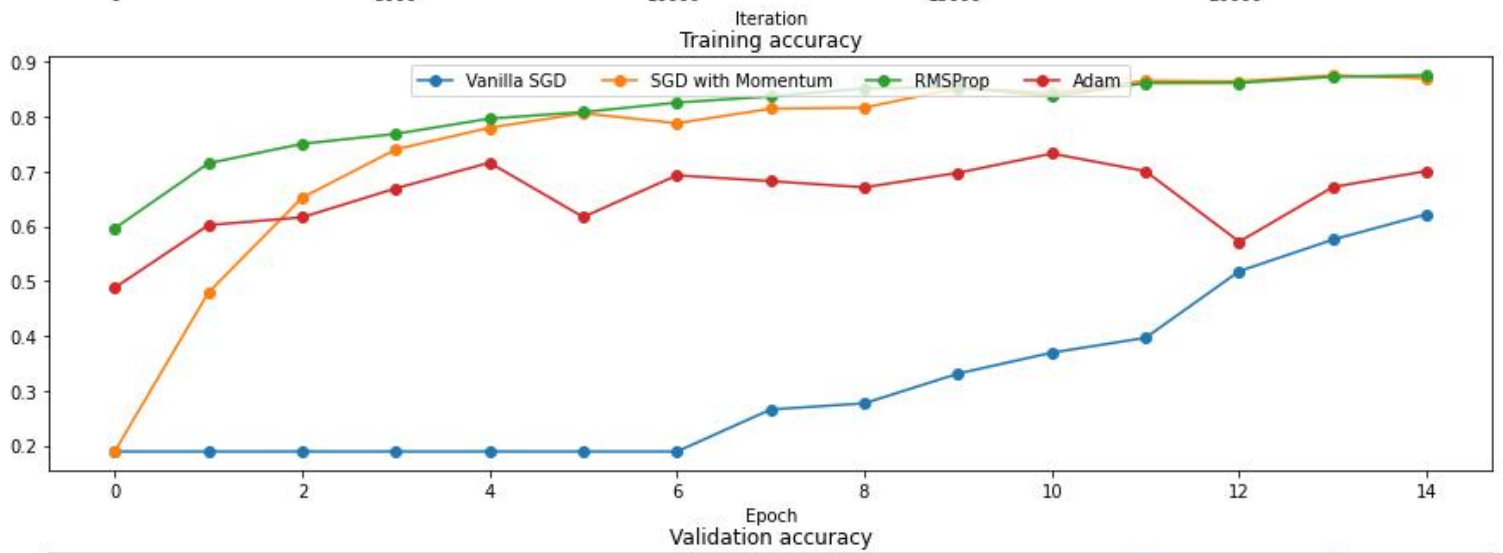

Training loss




Accuracy









Sigmoid




LeakyReLU




ReLU




ELU



SeLU



Tanh




```
[8]:
small_data_dict = {
    "data_train": data_train,
    "data_val": data_val,
    "data_test": (data_test, data_test_val)
}

model_sgd = FullyConnectedNetwork()
loss_entropy = PolyLossEntropy()
loss_cross_entropy = CrossEntropy()
optimizer_sgd = SGD(model_sgd.net, lr=0.001)

print("Training with Vanilla SGD, Validation Accuracy: 0.18206938900828984")
results_sgd = train_net(small_data_dict, model_sgd, loss_f_sgd, optimizer_sgd, batch_size=32,
                        max_epochs=15, show_every=100, verbose=True)

opt_params_sgd, loss_hist_sgd, train_acc_hist_sgd, val_acc_hist_sgd = results_sgd
opt_params_sgd, loss_hist_sgd, train_acc_hist_sgd, val_acc_hist_sgd = results_sgd

plt.subplot(3, 1, 1)
plt.title("Training loss")
plt.xlabel("Epoch")

plt.subplot(3, 1, 2)
plt.title("Training accuracy")
plt.xlabel("Epoch")

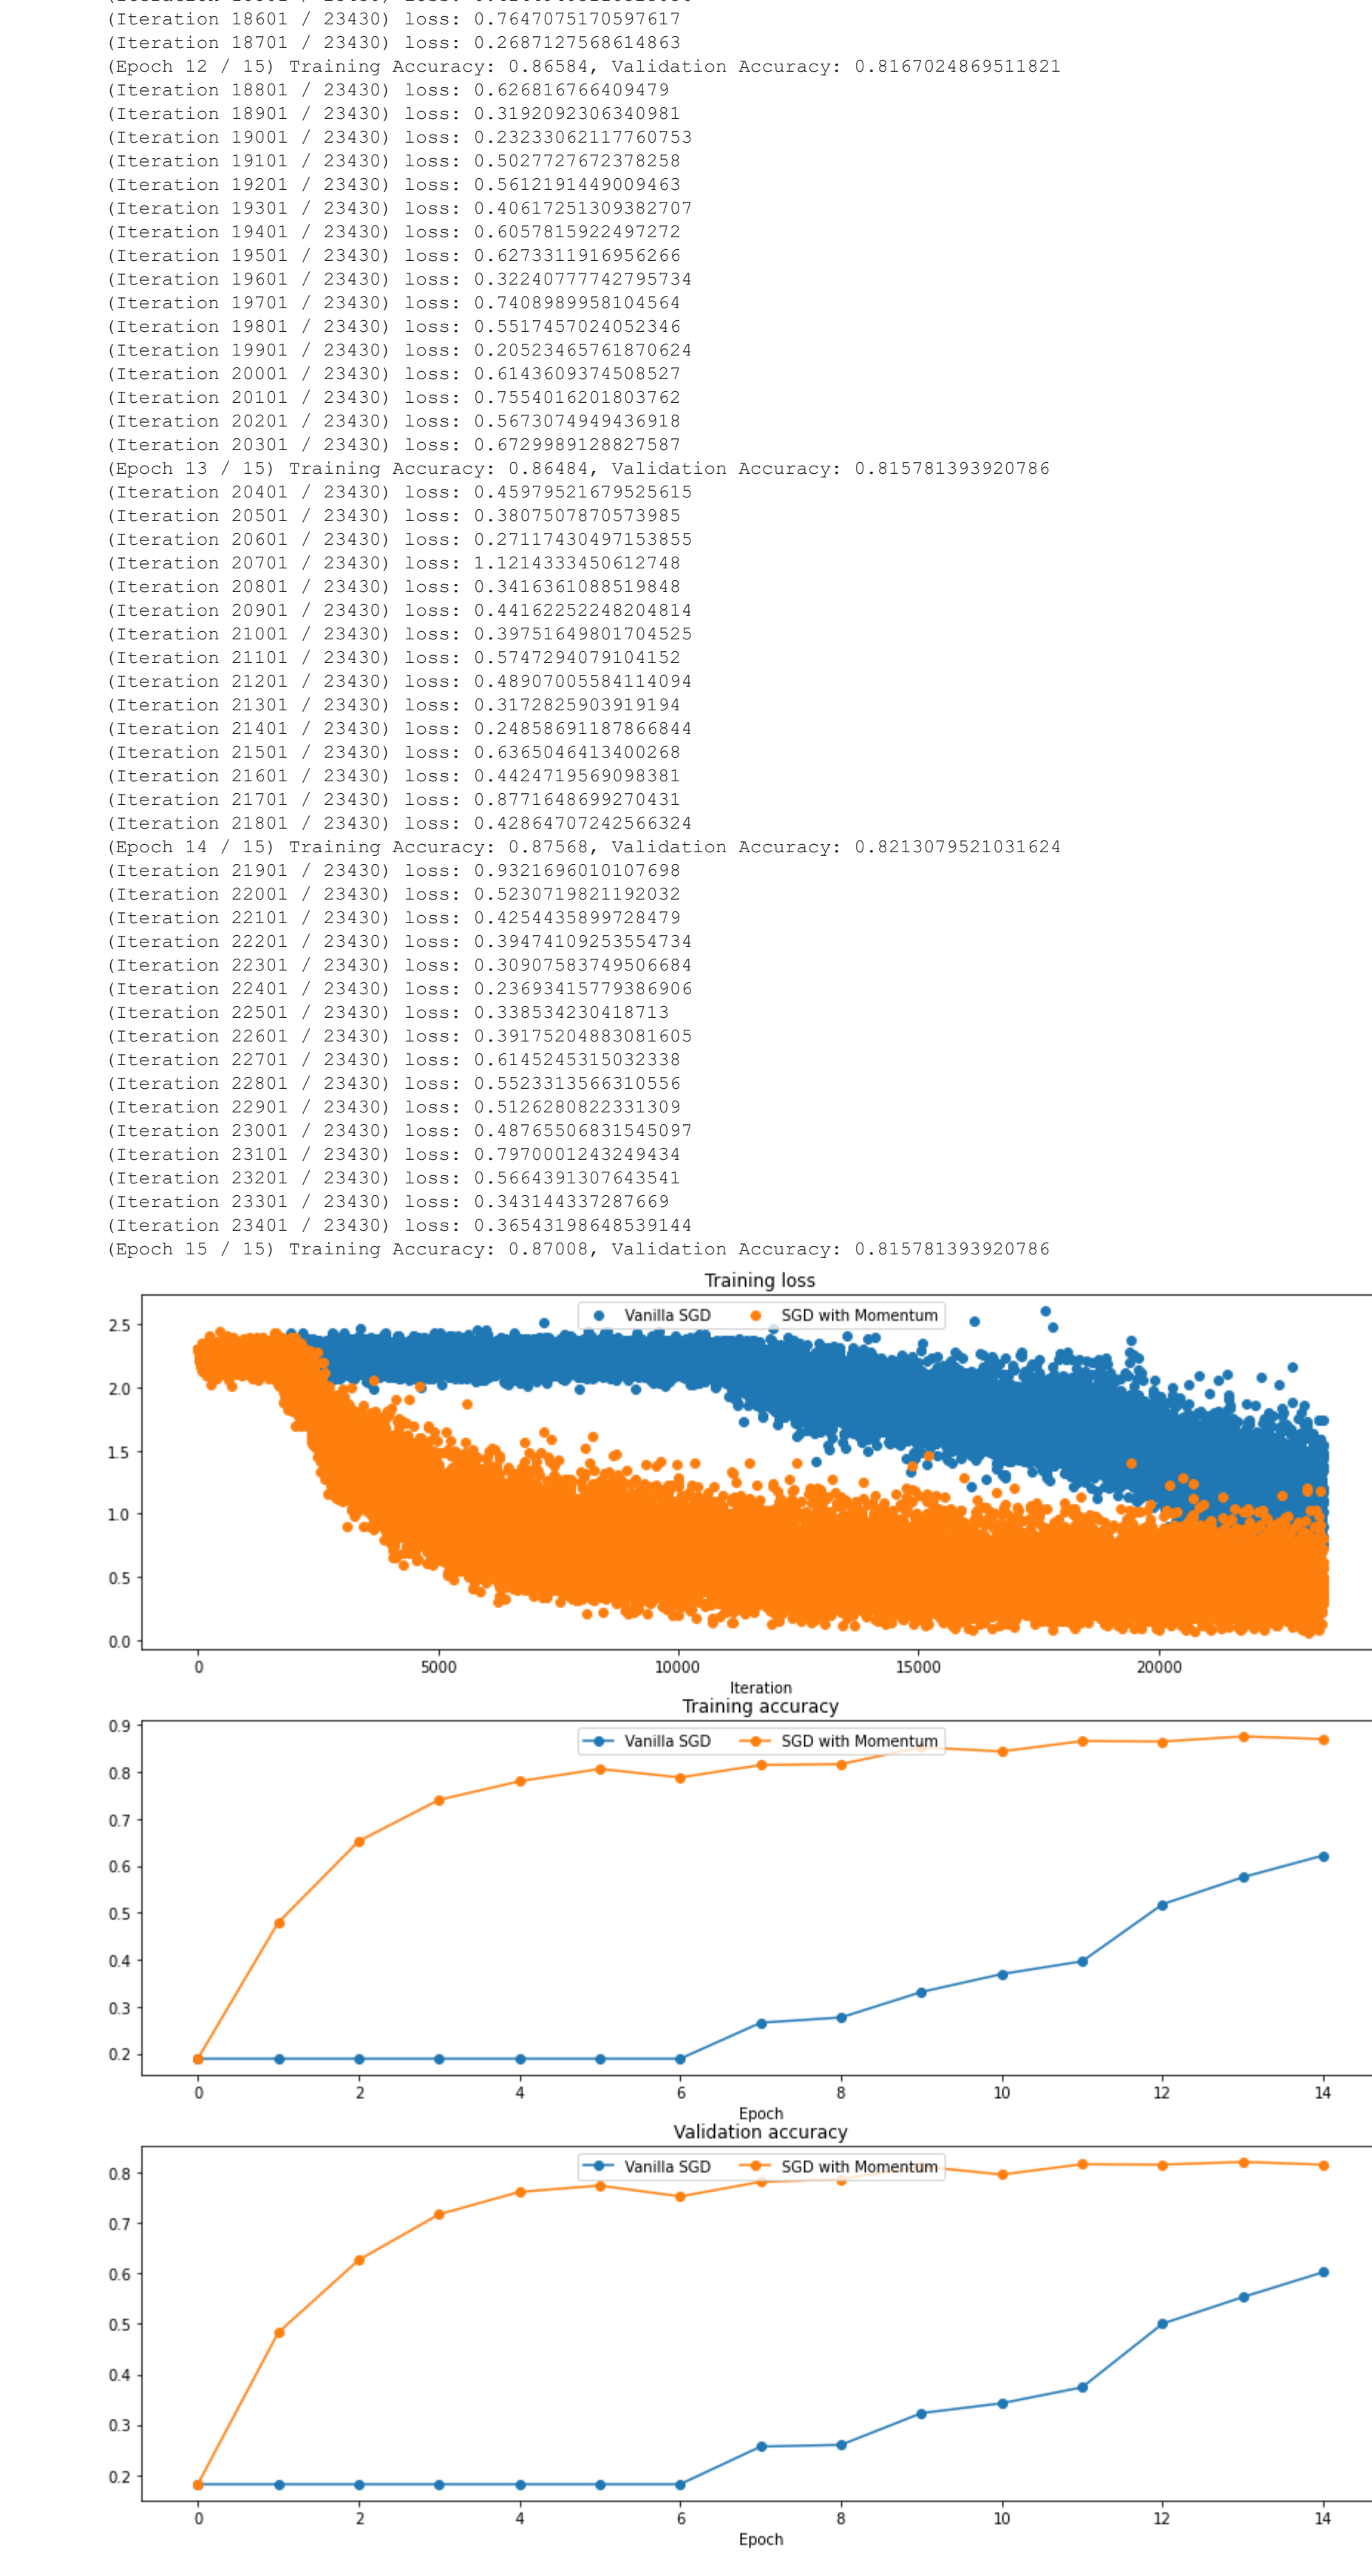
plt.subplot(3, 1, 3)
plt.title("Validation accuracy")
plt.xlabel("Epoch")

plt.subplot(3, 1, 1)
plt.plot(loss_hist_sgd, 'o', label="Vanilla SGD")
plt.subplot(3, 1, 2)
plt.plot(train_acc_hist_sgd, 'o', label="Vanilla SGD")
plt.subplot(3, 1, 3)
plt.plot(val_acc_hist_sgd, 'o', label="Vanilla SGD")

plt.subplot(3, 1, 1)
plt.plot(loss_hist_sgd, 'o', label="SGD with Momentum")
plt.subplot(3, 1, 2)
plt.plot(train_acc_hist_sgd, 'o', label="SGD with Momentum")
plt.subplot(3, 1, 3)
plt.plot(val_acc_hist_sgd, 'o', label="SGD with Momentum")

for i in [1, 2, 3]:
    plt.legend(loc=i)
plt.legend(loc="upper center", ncol=4)
plt.savefig('net_fig1.png', dpi=15)
plt.show()

Training with Vanilla SGD...
(iteration 1 / 2340) loss: 2.30202405405127
(iteration 101 / 2340) loss: 2.303440123828899
(iteration 201 / 2340) loss: 2.29836105262639
(iteration 301 / 2340) loss: 2.24375710272611
(iteration 401 / 2340) loss: 2.23831901052639
(iteration 501 / 2340) loss: 2.25748533373464
(iteration 601 / 2340) loss: 2.245168770884394
(iteration 701 / 2340) loss: 2.262214232938213
(iteration 801 / 2340) loss: 2.257379548640166
(iteration 901 / 2340) loss: 2.380941059089577
(iteration 1001 / 2340) loss: 2.20257836610278
(iteration 1101 / 2340) loss: 2.193849105641696
(iteration 1201 / 2340) loss: 2.142986797642235
(iteration 1301 / 2340) loss: 2.2668718233746474
(iteration 1401 / 2340) loss: 2.150475377731493
(iteration 1501 / 2340) loss: 2.273529705068405
Epoch 1 / 15 Training Accuracy: 0.1889, Validation Accuracy: 0.18206938900828984
(iteration 1601 / 2340) loss: 2.24601413303293
(iteration 1701 / 2340) loss: 2.333783093042975
(iteration 1801 / 2340) loss: 2.11751392081804
(iteration 1901 / 2340) loss: 2.205548519231465
(iteration 2001 / 2340) loss: 2.143731801898435
(iteration 2101 / 2340) loss: 2.2396229891697548
(iteration 2201 / 2340) loss: 2.235938509131465
(iteration 2301 / 2340) loss: 2.123998360305853
(iteration 2401 / 2340) loss: 2.257688844005743
(iteration 2501 / 2340) loss: 2.269723409642046
(iteration 2601 / 2340) loss: 2.21627896362204
(iteration 2701 / 2340) loss: 2.142043050466814
(iteration 2801 / 2340) loss: 2.239520105412478
(iteration 2901 / 2340) loss: 2.267620327979775
(iteration 3001 / 2340) loss: 2.309751637777322
(iteration 3101 / 2340) loss: 2.25653253220466
Epoch 2 / 15 Training Accuracy: 0.1889, Validation Accuracy: 0.18206938900828984
(iteration 3201 / 2340) loss: 2.194154112103643
(iteration 3301 / 2340) loss: 2.339520105412478
(iteration 3401 / 2340) loss: 2.215249933798637
(iteration 3501 / 2340) loss: 2.167523205841247
(iteration 3601 / 2340) loss: 2.189786138915265
(iteration 3701 / 2340) loss: 2.13126257298502
(iteration 3801 / 2340) loss: 2.249513115942046
(iteration 3901 / 2340) loss: 2.28545946809294
(iteration 4001 / 2340) loss: 2.198251521053874
(iteration 4101 / 2340) loss: 2.207460073142493
(iteration 4201 / 2340) loss: 2.22467412692337
(iteration 4301 / 2340) loss: 2.05303939609147
(iteration 4401 / 2340) loss: 2.35274699642046
(iteration 4501 / 2340) loss: 2.229357460968944
Epoch 3 / 15 Training Accuracy: 0.1889, Validation Accuracy: 0.18206938900828984
(iteration 4601 / 2340) loss: 2.397248473142493
(iteration 4701 / 2340) loss: 2.267153104703051
(iteration 4801 / 2340) loss: 2.083549105641696
(iteration 4901 / 2340) loss: 2.207460073142493
(iteration 5001 / 2340) loss: 2.181381060556312
(iteration 5101 / 2340) loss: 2.147264874797247
(iteration 5201 / 2340) loss: 2.17789056412478
(iteration 5301 / 2340) loss: 2.19372989103671
(iteration 5401 / 2340) loss: 2.139454956297456
(iteration 5501 / 2340) loss: 2.220054956297456
(iteration 5601 / 2340) loss: 2.2623237466631857
(iteration 5701 / 2340) loss: 2.22768557402541
(iteration 5801 / 2340) loss: 2.088270626262626
(iteration 5901 / 2340) loss: 2.185019494228804
(iteration 6001 / 2340) loss: 2.193125857353622
(iteration 6101 / 2340) loss: 2.273453520460166
Epoch 4 / 15 Training Accuracy: 0.1889, Validation Accuracy: 0.18206938900828984
(iteration 6201 / 2340) loss: 2.233184180123763
(iteration 6301 / 2340) loss: 2.170078382626262
(iteration 6401 / 2340) loss: 2.263505258663023
(iteration 6501 / 2340) loss: 2.119606956412478
(iteration 6601 / 2340) loss: 2.207306433804298
(iteration 6701 / 2340) loss: 2.3338415868150264
(iteration 6801 / 2340) loss: 2.2429241142493
(iteration 6901 / 2340) loss: 2.19814103372024
(iteration 7001 / 2340) loss: 2.261579850472472
(iteration 7101 / 2340) loss: 2.396177404601696
(iteration 7201 / 2340) loss: 2.263987570803724
(iteration 7301 / 2340) loss: 2.2073632924393647
(iteration 7401 / 2340) loss: 2.1837493861371
(iteration 7501 / 2340) loss: 2.27381979361371
(iteration 7601 / 2340) loss: 2.1826563818643277
Epoch 5 / 15 Training Accuracy: 0.1889, Validation Accuracy: 0.18206938900828984
(iteration 7701 / 2340) loss: 2.278704444444444
(iteration 7801 / 2340) loss: 2.361265166117029
(iteration 7901 / 2340) loss: 2.180538946323763
(iteration 8001 / 2340) loss: 2.29861489103671
(iteration 8101 / 2340) loss: 2.2363079019196635
(iteration 8201 / 2340) loss: 2.258491982887293
(iteration 8301 / 2340) loss: 2.241391982887293
(iteration 8401 / 2340) loss: 2.2115104713312093
(iteration 8501 / 2340) loss: 2.333687143253586
(iteration 8601 / 2340) loss: 2.146024978641696
(iteration 8701 / 2340) loss: 2.2739742890572265
(iteration 8801 / 2340) loss: 2.278883135029074
(iteration 8901 / 2340) loss: 2.101370977142493
(iteration 9001 / 2340) loss: 2.21753465825545
(iteration 9101 / 2340) loss: 2.276158675387075
Epoch 6 / 15 Training Accuracy: 0.1889, Validation Accuracy: 0.18206938900828984
(iteration 9201 / 2340) loss: 2.0642681789834
(iteration 9301 / 2340) loss: 2.2745094868034987
(iteration 9401 / 2340) loss: 2.133346288197944
(iteration 9501 / 2340)
```



The update rule of

$$E[\epsilon^2]$$

```

[gt]it = γE[gt]it-1 + (1 - γ)git
θt+1 = θt -
    √E[gt]it / (D + 1)
in lib/optim.py

```

```
opt_rms = RMSProp(test_rms, 1e-2, 0.9)
opt_rms.cache = {"rms_fc_w": cache}
```

```

opt_params.update(w = test_layers[0].params["rms_fc_w"])
cache = opt_rms.cache["rms_fc_w"]

expected_updated_w = np.asarray([
    -0.39232849, -0.34037513, -0.28849239, -0.23659121, -0.18467247,
    -0.132737, -0.08078555, -0.02863884, 0.02316247, 0.07513774,
    -0.12716641, 0.17918792, 0.23122175, 0.28326742, 0.33532447,
    0.38739248, 0.43947102, 0.49155973, 0.54365823, 0.59576619]])

expected_cache = np.asarray([
    0.5976, 0.61622777, 0.6277108, 0.64284931, 0.658084321,
    0.67329252, 0.68859723, 0.70395734, 0.71937285, 0.73478477,
    0.75037008, 0.76593518, 0.78158892, 0.79728144, 0.81302393,
    0.82882629, 0.84469141, 0.86060534, 0.87657507, 0.8926 ]])

print ('The following errors should be around or less than 1e-7')
print ('updated_w_error:', rel_error(expected_updated_w, updated_w))
print ('cache_error:', rel_error(expected_cache, opt_rms.cache["rms_fc_w"]))

The following errors should be around or less than 1e-7
updated_w_error: 9.502546229894295e-08
cache_error: 2.6477955807156126e-09
```

Adam [2pt]

The update rule of Adam is as shown below:

$$\begin{aligned}
 t &= t + 1 \\
 g_t &= \text{gradients at update step } t \\
 m_t &= \beta_1 m_{t-1} + (1 - \beta_1) g_t \\
 v_t &= \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \\
 \hat{m}_t &= m_t / (1 - \beta_1^t) \\
 \hat{v}_t &= v_t / (1 - \beta_2^t) \\
 \theta_{t+1} &= \theta_t - \frac{\eta \hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}
 \end{aligned}$$

Complete the `Adam()` function in `lib/optin.py`. Important Notes: 1) `t` must be updated before everything else 2) β_1^t is β_1 exponentiated to the t 'th power

```
In [87]: @xload_ext autoreload
seed = 123
np.random.seed(seed=seed)

# Test Adam implementation; you should see errors around 1e-7 or less
N, D = 4, 5
test_adam = sequential(fc(N, D, name="adam_fc"))

w = np.linspace(-0.4, 0.6, num=N*D).reshape(N, D)
m = np.linspace(-0.6, 0.4, num=N*D).reshape(N, D)
v = np.linspace(0.6, 0.9, num=N*D).reshape(N, D)
v = np.linspace(0.7, 0.5, num=N*D).reshape(N, D)

test_adam.params[0].params = ["adam_fc_w", w]
test_adam.layers[0].grads = ["adam_fc_w": dw]

opt_adam = Adam(test_adam, le=2, 0.9, 0.999, t=5)
opt_adam.mt = ["adam_fc_w": m]
opt_adam.vt = ["adam_fc_w": v]
opt_adam.step()

updated_w = test_adam.layers[0].params["adam_fc_w"]
mt = opt_adam.mt["adam_fc_w"]
vt = opt_adam.vt["adam_fc_w"]

expected_updated_w = np.asarray([
    -0.40094747, -0.34836187, -0.29577703, -0.24319299, -0.190609771,
    -0.1380274, -0.08544591, -0.03286534, 0.01971428, 0.0722929,
    -0.1487805, -0.1774702, 0.23002243, 0.28259667, 0.33516969,
    0.38741745, 0.44031188, 0.49288203, 0.54544852, 0.59801459]])
expected_v = np.asarray([
    0.60000000, 0.60000000, 0.60000000, 0.60000000, 0.60000000,
    0.60000000, 0.60000000, 0.60000000, 0.60000000, 0.60000000,
```

```
expected_m
[ 0.48,
[ 0.5773,
[ 0.6747,
[ 0.7721]
```

```
print ('The following erro
print ('updated_w error: '
```

```

    # be around or less than 1e-7
    if (expected_m, v1)
        error(expected_m, v1)
    # 7.9531431e-07
    6e-09
    1e-09

```

mizers [4pt]

compare the plotted results among all the above optimizers. Why is Adam better than Vanilla SGD optimizer.

```
print ("Training with RMSProp...")
results_rms = train_net(small_data_dict, model_rms, loss_f_rms, optimizer_rms, batch_size=32,
```

```
max_epochs=15, show
```

```

results_adam = train_net(small_data_dict, model_adam, optimizer_adam, batch_size=32,
                          max_epochs=15, show_every=100, verbose=True)

opt_params_rms, loss_hist_rms, train_acc_hist_rms, val_acc_hist_rms = results_rms
opt_params_adam, loss_hist_adam, train_acc_hist_adam, val_acc_hist_adam = results_adam

plt.subplot(3, 1, 1)
plt.title('Training loss')
plt.xlabel('Iteration')

plt.subplot(3, 1, 2)
plt.title('Training accuracy')
plt.xlabel('Epoch')

plt.subplot(3, 1, 3)
plt.title('Validation accuracy')
plt.xlabel('Epoch')

plt.subplot(3, 1, 1)
plt.plot(loss_hist_sgd, 'o', label='Vanilla SGD')
plt.subplot(3, 1, 2)
plt.plot(train_acc_hist_sgd, '-o', label='Vanilla SGD')
plt.subplot(3, 1, 3)
plt.plot(val_acc_hist_sgd, '-o', label='Vanilla SGD')

plt.subplot(3, 1, 1)
plt.plot(loss_hist_sgdm, 'o', label='SGD with Momentum')
plt.subplot(3, 1, 2)
plt.plot(train_acc_hist_sgdm, '-o', label='SGD with Momentum')
plt.subplot(3, 1, 3)
plt.plot(val_acc_hist_sgdm, '-o', label='SGD with Momentum')

plt.subplot(3, 1, 1)
plt.plot(loss_hist_rms, 'o', label='RMSProp')
plt.subplot(3, 1, 2)
plt.plot(train_acc_hist_rms, '-o', label='RMSProp')
plt.subplot(3, 1, 3)
plt.plot(val_acc_hist_rms, '-o', label='RMSProp')

plt.subplot(3, 1, 1)
plt.plot(loss_hist_adam, 'o', label='Adam')
plt.subplot(3, 1, 2)
plt.plot(train_acc_hist_adam, '-o', label='Adam')
plt.subplot(3, 1, 3)
plt.plot(val_acc_hist_adam, '-o', label='Adam')

for i in (1, 2, 3):
    plt.subplot(3, 1, i)
    plt.legend(loc='upper center', ncol=4)
    plt.gcf().set_size_inches(15, 15)
plt.show()

```


Training with RMSProp...

Iteration 101 / 23430 losses: 2.30203620485127
(iteration 201 / 23430) losses: 1.893342872033863
(iteration 301 / 23430) losses: 1.650792248851311
(iteration 401 / 23430) losses: 1.590375310840349
(iteration 501 / 23430) losses: 2.273858354003456
(iteration 601 / 23430) losses: 1.727650187187131
(iteration 701 / 23430) losses: 1.680880525185862
(iteration 801 / 23430) losses: 1.667282835389891
(iteration 901 / 23430) losses: 1.602854629803895
(iteration 1001 / 23430) losses: 1.363650100786184
(iteration 1101 / 23430) losses: 1.427338284882828
(iteration 1201 / 23430) losses: 1.272650162562527
(iteration 1301 / 23430) losses: 1.199111477115944
(iteration 1401 / 23430) losses: 1.183286782020383
(iteration 1501 / 23430) losses: 1.2028549062980835
(Epoch 1 / 15) Training Accuracy: 0.59828, Validation Accuracy: 0.578753454098864
(iteration 1601 / 23430) losses: 1.005253569734358
(iteration 1701 / 23430) losses: 0.279321863293317
(iteration 1801 / 23430) losses: 0.814351205490478
(iteration 1901 / 23430) losses: 0.913094345876718
(iteration 2001 / 23430) losses: 0.689749097965034
(iteration 2101 / 23430) losses: 1.097701520428926
(iteration 2201 / 23430) losses: 1.0259748091139227
(iteration 2301 / 23430) losses: 0.724924906250249
(iteration 2401 / 23430) losses: 1.051623779151632
(iteration 2501 / 23430) losses: 1.051785973243555
(iteration 2601 / 23430) losses: 0.744501381871311
(iteration 2701 / 23430) losses: 1.126757846821684
(iteration 2801 / 23430) losses: 0.801893116476154
(iteration 2901 / 23430) losses: 0.646023892526658
(iteration 3001 / 23430) losses: 1.2576112977471592
(iteration 3101 / 23430) losses: 1.108383950954089
(iteration 3201 / 23430) losses: 0.9704861862928803
(Epoch 2 / 15) Training Accuracy: 0.74688, Validation Accuracy: 0.690819773797025
(iteration 3301 / 23430) losses: 0.5854400561402141
(iteration 3401 / 23430) losses: 0.583814921436063
(iteration 3501 / 23430) losses: 0.702455027088514
(iteration 3601 / 23430) losses: 0.560373779542691
(iteration 3701 / 23430) losses: 0.689749097965034
(iteration 3801 / 23430) losses: 0.794064984734206
(iteration 3901 / 23430) losses: 1.079932640527922
(iteration 4001 / 23430) losses: 0.52493169252523
(iteration 4101 / 23430) losses: 0.697303569270543
(iteration 4201 / 23430) losses: 0.843021394370783
(iteration 4301 / 23430) losses: 0.52493169252523
(iteration 4401 / 23430) losses: 0.881303863106439
(iteration 4501 / 23430) losses: 1.0635697898722
(iteration 4601 / 23430) losses: 0.9704861862928803
(Epoch 3 / 15) Training Accuracy: 0.75066, Validation Accuracy: 0.7279705250320273
(iteration 4701 / 23430) losses: 0.492694179728893
(iteration 4801 / 23430) losses: 0.899270165230617
(iteration 4901 / 23430) losses: 0.947465973201247
(iteration 5001 / 23430) losses: 0.4428209301166264
(iteration 5101 / 23430) losses: 0.780435815971715
(iteration 5201 / 23430) losses: 0.886058356584739
(iteration 5301 / 23430) losses: 0.887627925142569
(iteration 5401 / 23430) losses: 0.47194205901811753
(iteration 5501 / 23430) losses: 0.3868406240250484
(iteration 5601 / 23430) losses: 0.702455027088514
(iteration 5701 / 23430) losses: 1.015653424461971
(iteration 5801 / 23430) losses: 0.534621426048133
(iteration 5901 / 23430) losses: 0.702455027088514
(iteration 6001 / 23430) losses: 0.434885916484747
(iteration 6101 / 23430) losses: 0.771581426057842
(Epoch 4 / 15) Training Accuracy: 0.74688, Validation Accuracy: 0.744243168560246
(iteration 6201 / 23430) losses: 0.9418478690367347
(iteration 6301 / 23430) losses: 0.538811493430454
(iteration 6401 / 23430) losses: 0.689749097965034
(iteration 6501 / 23430) losses: 0.703679645347742
(iteration 6601 / 23430) losses: 0.629064521506778
(iteration 6701 / 23430) losses: 0.646023892526658
(iteration 6801 / 23430) losses: 0.6660227413159094
(iteration 6901 / 23430) losses: 0.9170356814312054
(iteration 7001 / 23430) losses: 0.635302314771718
(iteration 7101 / 23430) losses: 0.7312045059158002
(iteration 7201 / 23430) losses: 0.7746887819006347
(iteration 7301 / 23430) losses: 0.61699334400933
(iteration 7401 / 23430) losses: 0.646023892526658
(iteration 7501 / 23430) losses: 0.7746887819006347
(iteration 7601 / 23430) losses: 0.61699334400933
(iteration 7701 / 23430) losses: 0.450495237990309
(iteration 7801 / 23430) losses: 0.450495237990309
(Epoch 5 / 15) Training Accuracy: 0.79666, Validation Accuracy: 0.7657353392692662
(iteration 7901 / 23430) losses: 0.569081787826217
(iteration 8001 / 23430) losses: 0.4510437967655743
(iteration 8101 / 23430) losses: 0.502778384084547
(iteration 8201 / 23430) losses: 0.452139075803732
(iteration 8301 / 23430) losses: 0.5677077210733411
(iteration 8401 / 23430) losses: 0.587598255520508
(iteration 8501 / 23430) losses: 0.724349387154169
(iteration 8601 / 23430) losses: 0.422472423035692
(iteration 8701 / 23430) losses: 0.888945027573546
(iteration 8801 / 23430) losses: 0.364473931521598
(iteration 8901 / 23430) losses: 0.763368617663068
(iteration 9001 / 23430) losses: 0.763368617663068
(iteration 9101 / 23430) losses: 0.41670133186450786
(iteration 9201 / 23430) losses: 0.6170356814312054
(iteration 9301 / 23430) losses: 0.795281610897778
(Epoch 6 / 15) Training Accuracy: 0.79666, Validation Accuracy: 0.768056493703964
(iteration 9401 / 23430) losses: 0.434885916484747
(iteration 9501 / 23430) losses: 0.622011743341585
(iteration 9601 / 23430) losses: 0.689749097965034
(iteration 9701 / 23430) losses: 0.3978822046925314
(iteration 9801 / 23430) losses: 0.366165541183911
(iteration 9901 / 23430) losses: 0.650431925865075
(iteration 10001 / 23430) losses: 0.865757042928939
(iteration 10101 / 23430) losses: 0.490649288064193
(iteration 10201 / 23430) losses: 0.089132781023063
(iteration 10301 / 23430) losses: 0.663105659281878
(iteration 10401 / 23430) losses: 0.47172273075927296
(iteration 10501 / 23430) losses: 0.509270165230617
(iteration 10601 / 23430) losses: 0.493521762672048
(iteration 10701 / 23430) losses: 0.698048513037085
(iteration 10801 / 23430) losses: 0.8639922977200148
(Epoch 7 / 15) Training Accuracy: 0.8259, Validation Accuracy: 0.7924470371507523
(iteration 10901 / 23430) losses: 0.2992450163401323
(iteration 11001 / 23430) losses: 0.4958313802605528
(iteration 11101 / 23430) losses: 0.8986486500236904
(iteration 11201 / 23430) losses: 0.52493169252523
(iteration 11301 / 23430) losses: 0.6111640272696471
(iteration 11401 / 23430) losses: 0.348387928728496
(iteration 11501 / 23430) losses: 0.188279050826168
(iteration 11601 / 23430) losses: 0.3631304060621907
(iteration 11701 / 23430) losses: 0.271215820693497
(iteration 11801 / 23430) losses: 0.402013503336983
(iteration 11901 / 23430) losses: 0.4264752167210515
(iteration 12001 / 23430) losses: 0.38392019974429253
(iteration 12101 / 23430) losses: 0.650431925865075
(iteration 12201 / 23430) losses: 0.37619369973909733
(iteration 12301 / 23430) losses: 0.5882029441391872
(Epoch 8 / 15) Training Accuracy: 0.8259, Validation Accuracy: 0.79606485108996
(iteration 12401 / 23430) losses: 0.4760155831301155
(iteration 12501 / 23430) losses: 0.4760155831301155
(iteration 12601 / 23430) losses: 0.3322468376752496
(iteration 12701 / 23430) losses: 0.58484902835703184
(iteration 12801 / 23430) losses: 0.658495406793798
(iteration 12901 / 23430) losses: 0.599711776389415
(iteration 13001 / 23430) losses: 0.658495406793798
(iteration 13101 / 23430) losses: 0.4737671970971019
(iteration 13201 / 23430) losses: 0.381130416041856
(iteration 13301 / 23430) losses: 0.46062918703783
(iteration 13401 / 23430) losses: 0.252701609504624
(iteration 13501 / 23430) losses: 0.4398866898448554
(iteration 13601 / 23430) losses: 0.4999886698448554
(iteration 13701 / 23430) losses: 0.653105659281878
(iteration 13801 / 23430) losses: 0.6328387896531356
(Epoch 9 / 15) Training Accuracy: 0.85176, Validation Accuracy: 0.7982806263423607
(iteration 13901 / 23430) losses: 0.2464501363061375
(iteration 14001 / 23430) losses: 0.17885846172671
(iteration 14101 / 23430) losses: 0.400241481729435
(iteration 14201 / 23430) losses: 0.400241481729435
(iteration 14301 / 23430) losses: 0.6836762329825
(iteration 14401 / 23430) losses: 0.418804681469459
(iteration 14501 / 23430) losses: 0.0702381682080618
(iteration 14601 / 23430) losses: 0.39802727880254
(iteration 14701 / 23430) losses: 0.2398404831366955
(iteration 14801 / 23430) losses: 0.0702381682080618
(iteration 14901 / 23430) losses: 0.3389692110631317
(iteration 15001 / 23430) losses: 0.519384649489884
(iteration 15101 / 23430) losses: 0.136682842926209
(iteration 15201 / 23430) losses: 0.44270261584156767
(iteration 15301 / 23430) losses: 0.500803907862059
(iteration 15401 / 23430) losses: 0.658495406793798
(Epoch 10 / 15) Training Accuracy: 0.85176, Validation Accuracy: 0.8094607737318156
(iteration 15501 / 23430) losses: 0.3752749503972613
(iteration 15601 / 23430) losses: 0.658495406793798
(iteration 15701 / 23430) losses: 0.4559657325077266
(iteration 15801 / 23430) losses: 0.480663846071934
(iteration 15901 / 23430) losses: 0.658495406793798
(iteration 16001 / 23430) losses: 0.6396420282123783
(iteration 16101 / 23430) losses: 0.550180584628028
(iteration 16201 / 23430) losses: 0.658495406793798
(iteration 16301 / 23430) losses: 0.156536940343688
(iteration 16401 / 23430) losses: 0.284478891789227
(iteration 16501 / 23430) losses: 0.509270165230617
(iteration 16601 / 23430) losses: 0.11640961348012807
(iteration 16701 / 23430) losses: 0.2615797020911103
(iteration 16801 / 23430) losses: 0.509270165230617
(iteration 16901 / 23430) losses: 0.5164441621543014
(Epoch 11 / 15) Training Accuracy: 0.87322, Validation Accuracy: 0.7875345409886388
(iteration 17001 / 23430) losses: 0.400241481729435
(iteration 17101 / 23430) losses: 0.593215208286972
(iteration 17201 / 23430) losses: 0.283833613883417
(iteration 17301 / 23430) losses: 0.493024906250249
(iteration 17401 / 23430) losses: 0.493024906250249
(iteration 17501 / 23430) losses: 0.509270165230617
(iteration 17601 / 23430) losses: 0.666071716094697
(iteration 17701 / 23430) losses: 0.389389682892177
(iteration 17801 / 23430) losses: 0.509270165230617
(iteration 17901 / 23430) losses: 0.666071716094697
(iteration 18001 / 23430) losses: 0.927959665231021
(iteration 18101 / 23430) losses: 0.4644552058834354
(iteration 18201 / 23430) losses: 0.593013465303063
(iteration 18301 / 23430) losses: 0.18441430557628959
(iteration 18401 / 23430) losses: 0.83167480444095
(iteration 18501 / 23430) losses: 0.658495406793798
(Epoch 12 / 15) Training Accuracy: 0.86158, Validation Accuracy: 0.808718606877495
(iteration 18601 / 23430) losses: 0.619304523765867
(iteration 18701 / 23430) losses: 0.584308593054033
(iteration 18801 / 23430) losses: 0.3657855032844483
(iteration 18901 / 23430) losses: 0.7789819358134661
(iteration 19001 / 23430) losses: 0.29864078014323
(iteration 19101 / 23430) losses: 0.502256401334464
(iteration 19201 / 23430) losses: 0.290697597800201
(iteration 19301 / 23430) losses: 0.165838951080084
(iteration 19401 / 23430) losses: 0.38326929781250874
(iteration 19501 / 23430) losses: 0.566502139630907
(iteration 19601 / 23430) losses: 0.658495406793798
(iteration 19701 / 23430) losses: 0.226922206890181
(iteration 19801 / 23430) losses: 0.5892859623827063
(iteration 19901 / 23430) losses: 0.2747693691081357
(iteration 20001 / 23430) losses: 0.2567487781676011
(iteration 20101 / 23430) losses: 0.608898319180801
(Epoch 13 / 15) Training Accuracy: 0.86158, Validation Accuracy: 0.803807184525637
(iteration 20201 / 23430) losses: 0.46661249626390217
(iteration 20301 / 23430) losses: 0.2328467362784007
(iteration 20401 / 23430) losses: 0.400241481729435
(iteration 20501 / 23430) losses: 0.6689142061139746
(iteration 20601 / 23430) losses: 0.303472884578959
(iteration 20701 / 23430) losses: 0.5516206610757932
(iteration 20801 / 23430) losses: 0.62913649031826
(iteration 20901 / 23430) losses: 0.260212330231082
(iteration 21001 / 23430) losses: 0.2484651080221082
(iteration 21101 / 23430) losses: 0.16967458103193
(iteration 21201 / 23430) losses: 0.400241481729435
(iteration 21301 / 23430) losses: 0.45418563006015
(iteration 21401 / 23430) losses: 0.357444211783971
(iteration 21501 / 23430) losses: 0.658495406793798
(Epoch 14 / 15) Training Accuracy: 0.87314, Validation Accuracy: 0.813939207859939
(iteration 21601 / 23430) losses: 0.26104652269647
(iteration 21701 / 23430) losses: 0.474436157004943
(iteration 21801 / 23430) losses: 0.3200118103049972
(iteration 21901 / 23430) losses: 0.4090921954398075
(iteration 22001 / 23430) losses: 0.584308593054033
(iteration 22101 / 23430) losses: 0.37746156973147194
(iteration 22201 / 23430) losses: 0.4676181243160977
(iteration 22301 / 23430) losses: 0.56317382831772
(iteration 22401 / 23430) losses: 0.264663958668935
(iteration 22501 / 23430) losses: 0.08584954030324
(iteration 22601 / 23430) losses: 0.414723635278175
(iteration 22701 / 23430) losses: 0.4331833603097
(iteration 22801 / 23430) losses: 0.5892859623827063
(iteration 22901 / 23430) losses: 0.6836915984166925
(iteration 23001 / 23430) losses: 0.5480110253489261
(Epoch 15 / 15) Training Accuracy: 0.87596, Validation Accuracy: 0.802886094592421

Training with Adam...

Iteration 101 / 23430 losses: 2.30213557093081
(iteration 201 / 23430) losses: 2.3397301907548504
(iteration 301 / 23430) losses: 2.218307762347379
(iteration 401 / 23430) losses: 1.954931850319041
(iteration 501 / 23430) losses: 2.2243776873397856
(iteration 601 / 23430) losses: 1.805812906894743
(iteration 701 / 23430) losses: 1.8505812906894743
(iteration 801 / 23430) losses: 2.080378677521474
(iteration 901 / 23430) losses: 1.98510240131222
(iteration 1001 / 23430) losses: 1.7001985645009372
(iteration 1101 / 23430) losses: 1.6577951491524003
(iteration 1201 / 23430) losses: 1.702431849637818
(iteration 1301 / 23430) losses: 1.5695030746651089
(iteration 1401 / 23430) losses: 1.4365129493884
(iteration 1501 / 23430) losses: 1.4561829493884
(Epoch 1 / 15) Training Accuracy: 0.4817316548971466
(iteration 1601 / 23430) losses: 1.496189956450157
(iteration 1701 / 23430) losses: 1.001548190320549
(iteration 1801 / 23430) losses: 1.213602897889361
(iteration 1901 / 23430) losses: 1.3502682174760737
(iteration 2001 / 23430) losses: 1.26443103244644
(iteration 2101 / 23430) losses: 1.213586340318645
(iteration 2201 / 23430) losses: 1.28091341910241
(iteration 2301 / 23430) losses: 1.419321650102312
(iteration 2401 / 23430) losses: 1.538047165203936
(iteration 2501 / 23430) losses: 1.3427795049612182
(iteration 2601 / 23430) losses: 1.414832866797919
(iteration 2701 / 23430) losses: 1.6597152844654828
(iteration 2801 / 23430) losses: 1.318874070990905
(iteration 2901 / 23430) losses: 1.6597152844654828
(iteration 3001 / 23430) losses: 1.036596236650542
(iteration 3101 / 23430) losses: 1.1659878628770064
(Epoch 2 / 15) Training Accuracy: 0.61022, Validation Accuracy: 0.593791735935332
(iteration 3201 / 23430) losses: 0.9973172749628125
(iteration 3301 / 23430) losses: 1.3582049246265029
(iteration 3401 / 23430) losses: 0.412397272361661
(iteration 3501 / 23430) losses: 1.2488820992083194
(iteration 3601 / 23430) losses: 1.638182041280537
(iteration 3701 / 23430) losses: 0.959369342696727
(iteration 3801 / 23430) losses: 1.023915448081233
(iteration 3901 / 23430) losses: 0.959369342696727
(iteration 4001 / 23430) losses: 1.316172734619392
(iteration 4101 / 23430) losses: 1.1208447817579227
(iteration 4201 / 23430) losses: 0.35693628470364
(iteration 4301 / 23430) losses: 1.279003180208353
(iteration 4401 / 23430) losses: 0.61048553913702
(iteration 4501 / 23430) losses: 1.046117928825
(Epoch 3 / 15) Training Accuracy: 0.67648, Validation Accuracy: 0.6014737488486337
(iteration 4601 / 23430) losses: 1.496189956450157
(iteration 4701 / 23430) losses: 0.439430913931496
(iteration 4801 / 23430) losses: 0.95429487263039
(iteration 4901 / 23430) losses: 1.05429487263039
(iteration 5001 / 23430) losses: 1.199803148110012
(iteration 5101 / 23430) losses: 0.862923524076055
(iteration 5201 / 23430) losses: 1.4163921731614733
(iteration 5301 / 23430) losses: 1.65044850525591
(iteration 5401 / 23430) losses: 1.489301546530615
(iteration 5501 / 23430) losses: 0.95185418509388
(iteration 5601 / 23430) losses: 1.144691337470804
(iteration 5701 / 23430) losses: 0.474136028687377
(iteration 5801 / 23430) losses: 1.1274270141394676
(iteration 5901 / 23430) losses: 1.497057821714896
(iteration 6001 / 23430) losses: 0.5892859623827063
(iteration 6101 / 23430) losses: 1.378488639643167
(iteration 6201 / 23430) losses: 1.292118564079035
(Epoch 4 / 15) Training Accuracy: 0.6932, Validation Accuracy: 0.6509357414798894
(iteration 6301 / 23430) losses: 0.908632088647546
(iteration 6401 / 23430) losses: 1.216398115123227

PyTorch 566 Deep Learning\HW2_2\cscse16-main\lib\mlp\layer_utils.py:360: RuntimeWarning: divide by zero encountered in log

loss = - np.sum(np.multiply(label, np.log(logits)) / feat.shape[0])

PyTorch 566 Deep Learning\HW2_2\cscse16-main\lib\mlp\layer_utils.py:360: RuntimeWarning: invalid va

loss = - np.sum(np.multiply(label, np.log(logits)) / feat.shape[0])

Loss encountered in log

Iteration 101 / 23430 losses: 1.746302971335626
(iteration 201 / 23430) losses: 1.364673853217598
(iteration 301 / 23430) losses: 1.629848547953378
(iteration 401 / 23430) losses: 1.629848547953378
(iteration 501 / 23430) losses: 0.995377522318354
(iteration 601 / 23430) losses: 1.163927817894036
(iteration 701 / 23430) losses: 0.954823336574356
(iteration 801 / 23430) losses: 1.064736