


## Question 1: AE-Sampling






**Inline Question: Describe your observations, why do you think they occur? [2pt]** \ (please limit your answer to <150 words) \ **Answer:**

The decoding model has learned the representation of the MNIST dataset in such a way that it can reconstruct the original image from this reduced dimensional space with minimal error. Here it is trying to reconstruct the random values to a digit, and we can see that it resembles a few digits, like 8 and 3. These are the representations it has learned.


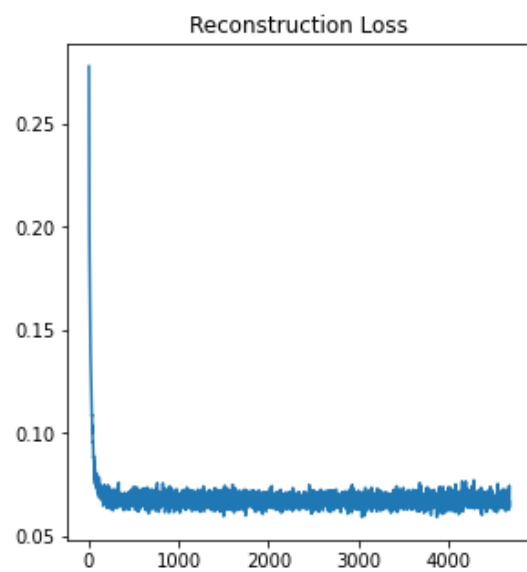
## Question 2:


B = 0:








B = 10:





B=0.0001 (opt):





Question 3: VAE inline questions

**Inline Question: What can you observe when setting  $\beta = 0$ ? Explain your observations! [3pt]** \ (please limit your answer to <150 words) \ **Answer:**

+ Code

+ Markdown

When  $\beta=0$ , we use only  $\log p(x|z)$  and ignore the prior divergence. We only use the reconstruction error. Hence we get the image shown above, without any disentanglement.

Let's repeat the same experiment for  $\beta = 10$ , a very high value for the coefficient. You can modify the  $\beta$  value in the cell above and rerun it (it is okay to overwrite the outputs of the previous experiment, but **make sure to copy the visualizations of training curves, reconstructions and samples for  $\beta = 0$  into your solution PDF** before deleting them).

**Inline Question: What can you observe when setting  $\beta = 10$ ? Explain your observations! [3pt]** \ (please limit your answer to <200 words) \ **Answer:**

When  $\beta=10$ , we have a stronger constraint over the latent bottleneck. This greatly limits the representation capacity of  $z$  and results in further disentanglement, which in turn, results in the above image with very poor reconstruction. Optimal value exists between 0 and 10.


Now we can start tuning the beta value to achieve a good result. First describe what a "good result" would look like (focus what you would expect for reconstructions and sample quality).

**Inline Question: Characterize what properties you would expect for reconstructions (1pt) and samples (2pt) of a well-tuned VAE! [3pt]** \ (please limit your answer to <200 words) \ **Answer:**


A well tuned B-VAE will learn the disentangled representations while still having less reconstruction errors. They will however have lesser interpretability of the latent space.

#### Question 4: Interpolation images


1-7:



3-5:



6-9:



## Question 5: Interpolation inline

Repeat the experiment for different start / end labels and different samples. Describe your observations.

**Inline Question: Repeat the interpolation experiment with different start / end labels and multiple samples. Describe your observations! Focus on: \**


1. **How do AE and VAE embedding space interpolations differ? \**
2. **How do you expect these differences to affect the usefulness of the learned representation for downstream learning? \** (please limit your answer to <300 words)

**Answer:**

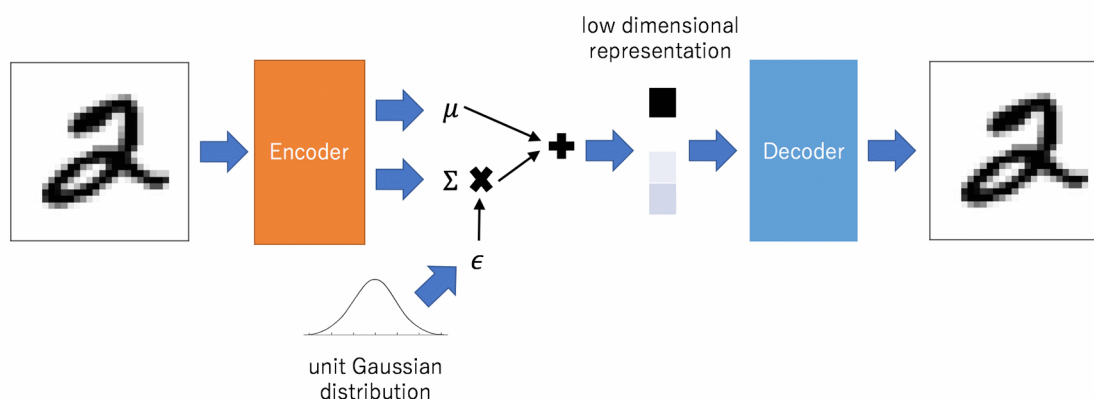
1. The latent space where the encoder encodes the input in an autoencoder may not be continuous and hence dont allow easy interpolation. On the other hand, the latent spaces of VAE are continous and allow random sampling.
2. VAE learns both the mean and variance of the inputs in the latent space and hence can be used to generate new data from feeding random values to the decoder part. The autoencoder will try its best to generate an image which looks close to a digit, while the VAE can generate an image closer to a digit because it learns a distribution of possible latent inputs that can lead to the specific digit.

Hence, we can observe that the transition between the 2 classes is smoother in the case of VAE and very abrupt (you can clearly see 2 different digits in the case of 3&5) in the case of AE. This is due to the nature of the latent space.

## Problem 1 - Variational Auto-Encoder (VAE)



Variational Auto-Encoders (VAEs) are a widely used class of generative models. They are simple to implement and, in contrast to other generative model classes like Generative Adversarial Networks, they optimize an explicit maximum likelihood objective to train the model. Finally, their architecture makes them well-suited for unsupervised representation learning, i.e. learning low-dimensional representations of high-dimensional inputs, like images, with only self-supervised objectives (data reconstruction in the case of VAEs).



(image source: <https://mlexplained.com/2017/12/28/an-intuitive-explanation-of-variational-autoencoders-vaes-part-1>)

**By working on this problem you will learn and practice the following steps:**

1. Set up a data loading pipeline in PyTorch.
2. Implement, train and visualize an auto-encoder architecture.
3. Extend your implementation to a variational auto-encoder.
4. Learn how to tune the critical beta parameter of your VAE.
5. Inspect the learned representation of your VAE.

**Note:** For faster training of the models in this assignment you can use Colab with enabled GPU support. In Colab, navigate to "Runtime" --> "Change Runtime Type" and set the "Hardware Accelerator" to "GPU".

### 1. MNIST Dataset

We will perform all experiments for this problem using the [MNIST dataset](#), a standard dataset of handwritten digits. The main benefits of this dataset are that it is small and relatively easy to model. It therefore allows for quick experimentation and serves as initial test bed in many papers.



Another benefit is that it is so widely used that PyTorch even provides functionality to automatically download it.

Let's start by downloading the data and visualizing some samples.

```
import matplotlib.pyplot as plt
%matplotlib inline

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

import torch
import torchvision

# this will automatically download the MNIST training set
mnist_train = torchvision.datasets.MNIST(root='./data',
                                         train=True,
                                         download=True,

transform=torchvision.transforms.ToTensor())
print("\n Download complete! Downloaded {} training
examples!".format(len(mnist_train)))

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-
ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-
ubyte.gz to ./data/MNIST/raw/train-images-idx3-ubyte.gz
{"version_major":2,"version_minor":0,"model_id":"f9a8cedde07e4b39a6f1a
81d5c26a87d"}

Extracting ./data/MNIST/raw/train-images-idx3-ubyte.gz to
./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-
ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-
ubyte.gz to ./data/MNIST/raw/train-labels-idx1-ubyte.gz
{"version_major":2,"version_minor":0,"model_id":"b6cf64adc5e0467ca658b
6985b299629"}

Extracting ./data/MNIST/raw/train-labels-idx1-ubyte.gz to
./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
to ./data/MNIST/raw/t10k-images-idx3-ubyte.gz
```

```
{"version_major":2,"version_minor":0,"model_id":"6c39f62de6fb457dbf34272df7cefd2b"}
```

Extracting ./data/MNIST/raw/t10k-images-idx3-ubyte.gz to  
./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz  
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz  
to ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz

```
{"version_major":2,"version_minor":0,"model_id":"f8e33690c8174c998aac65af0d34a287"}
```


Extracting ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz to  
./data/MNIST/raw

Download complete! Downloaded 60000 training examples!

```
/opt/conda/lib/python3.7/site-packages/torchvision/datasets/  
mnist.py:498: UserWarning: The given NumPy array is not writeable, and  
PyTorch does not support non-writeable tensors. This means you can  
write to the underlying (supposedly non-writeable) NumPy array using  
the tensor. You may want to copy the array to protect its data or make  
it writeable before converting it to a tensor. This type of warning  
will be suppressed for the rest of this program. (Triggered internally  
at /usr/local/src/pytorch/torch/csrc/utils/tensor_numpy.cpp:174.)  
    return torch.from_numpy(parsed.astype(m[2], copy=False)).view(*s)
```

```
import matplotlib.pyplot as plt  
import numpy as np
```

```
# Let's display some of the training samples.  
sample_images = []  
mnist_it = iter(mnist_train) # create simple iterator, later we will  
use proper DataLoader  
for _ in range(5):  
    sample = next(mnist_it) # samples a tuple (image, label)  
    sample_images.append(sample[0][0].data.cpu().numpy())  
  
fig = plt.figure(figsize = (10, 50))  
ax1 = plt.subplot(111)  
ax1.imshow(np.concatenate(sample_images, axis=1), cmap='gray')  
plt.show()
```



## 2. Auto-Encoder

Before implementing the full VAE, we will first implement an **auto-encoder architecture**. Auto-encoders feature the same encoder-decoder architecture as VAEs and therefore also learn a low-dimensional representation of the input data without supervision. In contrast to VAEs they are **fully deterministic** models and do not employ variational inference for optimization.

The **architecture** is very simple: we will encode the input image into a low-dimensional representation using a convolutional network with strided convolutions that reduce the image resolution in every layer. This results in a low-dimensional representation of the input image. This representation will get decoded back into the dimensionality of the input image using a convolutional decoder network that mirrors the architecture of the encoder. It employs transposed convolutions to increase the resolution of its input in every layer. The whole model is trained by **minimizing a reconstruction loss** between the input and the decoded image.

Intuitively, the **auto-encoder needs to compress the information contained in the input image** into a much lower dimensional representation (e.g.  $28 \times 28 = 784$ px vs. 64 embedding dimensions for our MNIST model). This is possible since the information captured in the pixels is *highly redundant*. E.g. encoding an MNIST image requires <4 bits to encode which of the 10 possible digits is displayed and a few additional bits to capture information about shape and orientation. This is much less than the  $255^{28 \cdot 28}$  bits of information that could be theoretically captured in the input image.

Learning such a **compressed representation can make downstream task learning easier**. For example, learning to add two numbers based on the inferred digits is much easier than performing the task based on two piles of pixel values that depict the digits.

In the following, we will first define the architecture of encoder and decoder and then train the auto-encoder model.

### Defining the Auto-Encoder Architecture [6pt]

```
import torch.nn as nn
```

```
# Let's define encoder and decoder networks
```

```
#####  
# Encoder Architecture:                                     #
```

```

# - Conv2d, hidden units: 32, output resolution: 14x14, kernel: 4 #
# - LeakyReLU #
# - Conv2d, hidden units: 64, output resolution: 7x7, kernel: 4 #
# - BatchNorm2d #
# - LeakyReLU #
# - Conv2d, hidden units: 128, output resolution: 3x3, kernel: 3 #
# - BatchNorm2d #
# - LeakyReLU #
# - Conv2d, hidden units: 256, output resolution: 1x1, kernel: 3 #
# - BatchNorm2d #
# - LeakyReLU #
# - Flatten #
# - Linear, output units: nz (= representation dimensionality) #
#####

```

```

class Encoder(nn.Module):
    def __init__(self, nz):
        super().__init__()
        ##### TODO
        #####
        # Create the network architecture using a nn.Sequential module
        # wrapper. #
        # All convolutional layers should also learn a bias.
        #
        # HINT: use the given information to compute stride and padding
        #
        # for each convolutional layer. Verify the shapes of
        # intermediate layers #
        # by running partial networks (with the next cell) and
        # visualizing the #
        # output shapes.
        #
        #####
        #####
        self.net = nn.Sequential(
            nn.Conv2d(1, 32, 4, stride=2, padding=1, bias=True),
            nn.LeakyReLU(),
            nn.Conv2d(32, 64, 4, stride=2, padding=1, bias=True),
            nn.BatchNorm2d(64),
            nn.LeakyReLU(),
            nn.Conv2d(64, 128, 3, stride=5, padding=1, bias=True),
            nn.BatchNorm2d(128),
            nn.LeakyReLU(),
            nn.Conv2d(128, 256, 3, stride=5, padding=1, bias=True),
            nn.BatchNorm2d(256),
            nn.LeakyReLU(),
            nn.Flatten(),
            nn.Linear(256, nz)
            # add your network layers here

```

```

        # ...
    )
    ##### END TODO
#####

def forward(self, x):
    return self.net(x)

#####
# Decoder Architecture (mirrors encoder architecture):
# - Linear, output units: 256
# - Reshape, output shape: (256, 1, 1)
# - BatchNorm2d
# - LeakyReLU
# - ConvT2d, hidden units: 128, output resolution: 3x3, kernel: 3
# - BatchNorm2d
# - LeakyReLU
# - ConvT2d, hidden units: 64, output resolution: 7x7, kernel: 3
# - ...
# - ...
# - ConvT2d, output units: 1, output resolution: 28x28, kernel: 4
# - Sigmoid (to limit output in range [0...1])
#####

class Decoder(nn.Module):
    def __init__(self, nz):
        super().__init__()
        ##### TODO
        #####
        # Create the network architecture using a nn.Sequential module
        # wrapper.
        # Again, all (transposed) convolutional layers should also learn a
        # bias.
        # We need to separate the initial linear layer into a separate
        # variable since
        # nn.Sequential does not support reshaping. Instead the "Reshape"
        # is performed
        # in the forward() function below and does not need to be added to
        # self.net
        # HINT: use the class nn.ConvTranspose2d for the transposed
        # convolutions.
        # Verify the shapes of intermediate layers by running
        # partial networks
        # (using the next cell) and visualizing the output shapes.
        #

#####
#####
self.map = nn.Linear(64, 256) # for initial Linear layer

```

```

self.net = nn.Sequential(
    nn.BatchNorm2d(256),
    nn.LeakyReLU(),
    nn.ConvTranspose2d(256, 128, 3, bias=True),
    nn.BatchNorm2d(128),
    nn.LeakyReLU(),
    nn.ConvTranspose2d(128, 64, 3, stride=2, padding=0,
bias=True),
    nn.BatchNorm2d(64),
    nn.LeakyReLU(),
    nn.ConvTranspose2d(64, 1, 4, stride=4, padding=0, bias=True),
    nn.Sigmoid()
    # add your network layers here
    # ...
)
##### END TODO
#####

def forward(self, x):
    return self.net(self.map(x).reshape(-1, 256, 1, 1))

```

### Testing the Auto-Encoder Forward Pass [1pt]

*# To test your encoder/decoder, let's encode/decode some sample images  
# first, make a PyTorch DataLoader object to sample data batches*

*batch\_size = 64*

*nworkers = 4 # number of wrokers used for efficient data loading*

```

##### TODO
#####
# Create a PyTorch DataLoader object for efficiently generating
training batches. #
# Make sure that the data loader automatically shuffles the training
dataset. #
# HINT: The DataLoader wraps the MNIST dataset class we created
earlier. #
# Use the given batch_size and number of data loading workers
when creating #
# the DataLoader.
#
#####
#####
mnist_data_loader = torch.utils.data.DataLoader(mnist_train,
batch_size=batch_size, shuffle=True, num_workers=nworkers)
##### END TODO
#####

```

*# now we can run a forward pass for encoder and decoder and check the produced shapes*

*nz = 64 # dimensionality of the learned embedding*

```

encoder = Encoder(nz)
decoder = Decoder(nz)
for sample_img, sample_label in mnist_data_loader:
    print(sample_img.shape)
    enc = encoder(sample_img)
    print("Shape of encoding vector (should be [batch_size, nz]):
    {}".format(enc.shape))
    dec = decoder(enc)
    print("Shape of decoded image (should be [batch_size, 1, 28, 28]):
    {}".format(dec.shape))
    break

```

/opt/conda/lib/python3.7/site-packages/torch/utils/data/dataloader.py:481: UserWarning: This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.

```

torch.Size([64, 1, 28, 28])
Shape of encoding vector (should be [batch_size, nz]): torch.Size([64, 64])
Shape of decoded image (should be [batch_size, 1, 28, 28]):
torch.Size([64, 1, 28, 28])

```

Now that we defined encoder and decoder network our architecture is nearly complete. However, before we start training, we can wrap encoder and decoder into an auto-encoder class for easier handling.

```

class AutoEncoder(nn.Module):
    def __init__(self, nz):
        super().__init__()
        self.encoder = Encoder(nz)
        self.decoder = Decoder(nz)

    def forward(self, x):
        return self.decoder(self.encoder(x))

    def reconstruct(self, x):
        """Only used later for visualization."""
        return self.forward(x)

```

## Setting up the Auto-Encoder Training Loop [6pt]

After implementing the network architecture, we can now set up the training loop and run training.

```

epochs = 10
learning_rate = 1e-3

```

```

# build AE model
device = torch.device('cuda:0' if torch.cuda.is_available() else
'cpu') # use GPU if available
ae_model = AutoEncoder(nz).to(device) # transfer model to GPU if
available
ae_model = ae_model.train() # set model in train mode (eg batchnorm
params get updated)

# build optimizer and loss function
##### TODO
#####
# Create the optimizer and loss classes. For the loss you can use a
loss layer #
# from the torch.nn package.
#
# HINT: We will use the Adam optimizer (learning rate given above,
otherwise #
# default parameters) and MSE loss for the criterion / loss.
#
# NOTE: We could also use alternative loss functions like cross
entropy, depending #
# on the assumptions we are making about the output
distribution. Here we #
# will use MSE loss as it is the most common choice, assuming a
Gaussian #
# output distribution.
#
#####
#####
opt = torch.optim.Adam(ae_model.parameters(), lr=learning_rate)
# create optimizer instance
criterion = nn.MSELoss() # create loss layer instance
##### END TODO
#####

train_it = 0
for ep in range(epochs):
    print("Run Epoch {}".format(ep))
    ##### TODO
    #####
    # Implement the main training loop for the auto-encoder model.
    #
    # HINT: Your training loop should sample batches from the data
    loader, run the #
    # forward pass of the AE, compute the loss, perform the
    backward pass and #
    # perform one gradient step with the optimizer.
    #
    # HINT: Don't forget to erase old gradients before performing the

```



```

backward pass.    #

#####
#####
# add training loop commands here
# ...
for (x, _) in mnist_data_loader:
    x = x.to(device)
    pred = ae_model.forward(x)
    rec_loss = criterion(pred, x)
    ae_model.zero_grad()
    opt.zero_grad()
    rec_loss.backward()
    opt.step()
##### END TODO
#####

    if train_it % 100 == 0:
        print("It {}: Reconstruction Loss: {}".format(train_it,
rec_loss))
        train_it += 1

print("Done!")

```

Run Epoch 0

```

It 0: Reconstruction Loss: 0.26972419023513794
It 100: Reconstruction Loss: 0.03894268348813057
It 200: Reconstruction Loss: 0.019504092633724213
It 300: Reconstruction Loss: 0.01590818539261818
It 400: Reconstruction Loss: 0.013952402397990227
It 500: Reconstruction Loss: 0.013186920434236526
It 600: Reconstruction Loss: 0.011852686293423176
It 700: Reconstruction Loss: 0.01009414717555046
It 800: Reconstruction Loss: 0.011076966300606728
It 900: Reconstruction Loss: 0.010262689553201199

```

Run Epoch 1

```

It 1000: Reconstruction Loss: 0.009367209859192371
It 1100: Reconstruction Loss: 0.00891298707574606
It 1200: Reconstruction Loss: 0.008198852650821209
It 1300: Reconstruction Loss: 0.00843334011733532
It 1400: Reconstruction Loss: 0.007510135415941477
It 1500: Reconstruction Loss: 0.008696877397596836
It 1600: Reconstruction Loss: 0.008465096354484558
It 1700: Reconstruction Loss: 0.0077793169766664505
It 1800: Reconstruction Loss: 0.008721661753952503

```

Run Epoch 2

```

It 1900: Reconstruction Loss: 0.007744807284325361
It 2000: Reconstruction Loss: 0.007225108332931995
It 2100: Reconstruction Loss: 0.007337048649787903
It 2200: Reconstruction Loss: 0.007528262212872505

```

It 2300: Reconstruction Loss: 0.006834432482719421  
It 2400: Reconstruction Loss: 0.007883540354669094  
It 2500: Reconstruction Loss: 0.00792559701949358  
It 2600: Reconstruction Loss: 0.0069623845629394054  
It 2700: Reconstruction Loss: 0.006955763790756464  
It 2800: Reconstruction Loss: 0.0073212794959545135

Run Epoch 3

It 2900: Reconstruction Loss: 0.00676879333332181  
It 3000: Reconstruction Loss: 0.006619803607463837  
It 3100: Reconstruction Loss: 0.006619736552238464  
It 3200: Reconstruction Loss: 0.006932450458407402  
It 3300: Reconstruction Loss: 0.006519296672195196  
It 3400: Reconstruction Loss: 0.006259030196815729  
It 3500: Reconstruction Loss: 0.0069648632779717445  
It 3600: Reconstruction Loss: 0.006809859536588192  
It 3700: Reconstruction Loss: 0.006310960277915001

Run Epoch 4

It 3800: Reconstruction Loss: 0.00664236955344677  
It 3900: Reconstruction Loss: 0.006767441052943468  
It 4000: Reconstruction Loss: 0.005658551584929228  
It 4100: Reconstruction Loss: 0.006615108344703913  
It 4200: Reconstruction Loss: 0.005319413263350725  
It 4300: Reconstruction Loss: 0.006962804589420557  
It 4400: Reconstruction Loss: 0.006364195607602596  
It 4500: Reconstruction Loss: 0.006309201940894127  
It 4600: Reconstruction Loss: 0.005965700838714838

Run Epoch 5

It 4700: Reconstruction Loss: 0.006194955203682184  
It 4800: Reconstruction Loss: 0.005932510830461979  
It 4900: Reconstruction Loss: 0.00668210769072175  
It 5000: Reconstruction Loss: 0.005948406644165516  
It 5100: Reconstruction Loss: 0.00530467601493001  
It 5200: Reconstruction Loss: 0.006765497848391533  
It 5300: Reconstruction Loss: 0.006568537559360266  
It 5400: Reconstruction Loss: 0.005505607929080725  
It 5500: Reconstruction Loss: 0.006162336561828852  
It 5600: Reconstruction Loss: 0.005985570140182972

Run Epoch 6

It 5700: Reconstruction Loss: 0.005671333055943251  
It 5800: Reconstruction Loss: 0.00556244095787406  
It 5900: Reconstruction Loss: 0.006094952579587698  
It 6000: Reconstruction Loss: 0.0058354646898806095  
It 6100: Reconstruction Loss: 0.005263315048068762  
It 6200: Reconstruction Loss: 0.006222279742360115  
It 6300: Reconstruction Loss: 0.005979042965918779  
It 6400: Reconstruction Loss: 0.005646682344377041  
It 6500: Reconstruction Loss: 0.0062297373078763485

Run Epoch 7

It 6600: Reconstruction Loss: 0.005251958500593901  
It 6700: Reconstruction Loss: 0.005530822090804577

```

It 6800: Reconstruction Loss: 0.00569359865039587
It 6900: Reconstruction Loss: 0.005079316440969706
It 7000: Reconstruction Loss: 0.0056202104315161705
It 7100: Reconstruction Loss: 0.005669177509844303
It 7200: Reconstruction Loss: 0.0064501347951591015
It 7300: Reconstruction Loss: 0.005197315476834774
It 7400: Reconstruction Loss: 0.005425662267953157
It 7500: Reconstruction Loss: 0.005457249004393816
Run Epoch 8
It 7600: Reconstruction Loss: 0.0051828231662511826
It 7700: Reconstruction Loss: 0.005223572254180908
It 7800: Reconstruction Loss: 0.005383913405239582
It 7900: Reconstruction Loss: 0.004581862594932318
It 8000: Reconstruction Loss: 0.005069257691502571
It 8100: Reconstruction Loss: 0.005365584511309862
It 8200: Reconstruction Loss: 0.005165536887943745
It 8300: Reconstruction Loss: 0.00555442413315177
It 8400: Reconstruction Loss: 0.005510886199772358
Run Epoch 9
It 8500: Reconstruction Loss: 0.005067567806690931
It 8600: Reconstruction Loss: 0.005851340480148792
It 8700: Reconstruction Loss: 0.004378402605652809
It 8800: Reconstruction Loss: 0.004706318955868483
It 8900: Reconstruction Loss: 0.005418729968369007
It 9000: Reconstruction Loss: 0.004921534564346075
It 9100: Reconstruction Loss: 0.005481590982526541
It 9200: Reconstruction Loss: 0.004407491069287062
It 9300: Reconstruction Loss: 0.0056149763986468315
Done!

```

## Verifying reconstructions [Opt]

Now that we trained the auto-encoder we can visualize some of the reconstructions on the test set to verify that it is converged and did not overfit. **Before continuing, make sure that your auto-encoder is able to reconstruct these samples near-perfectly.**

```

# visualize test data reconstructions
def vis_reconstruction(model):
    # download MNIST test set + build Dataset object
    mnist_test = torchvision.datasets.MNIST(root='./data',
                                             train=False,
                                             download=True,

transform=torchvision.transforms.ToTensor())
    mnist_test_iter = iter(mnist_test)
    model.eval()          # set model in evaluation mode (eg freeze
batchnorm params)
    input_imgs, test_reconstructions = [], []
    for _ in range(5):
        input_img = np.asarray(next(mnist_test_iter)[0])

```

```

        reconstruction = model.reconstruct(torch.tensor(input_img[None],
device=device))
        input_imgs.append(input_img[0])
        test_reconstructions.append(reconstruction[0,
0].data.cpu().numpy())

    fig = plt.figure(figsize = (20, 50))
    ax1 = plt.subplot(111)
    ax1.imshow(np.concatenate([np.concatenate(input_imgs, axis=1),
                                np.concatenate(test_reconstructions,
axis=1)], axis=0), cmap='gray')
    plt.show()

vis_reconstruction(ae_model)

```



## Sampling from the Auto-Encoder [2pt]

To test whether the auto-encoder is useful as a generative model, we can use it like any other generative model: draw embedding samples from a prior distribution and decode them through the decoder network. We will choose a unit Gaussian prior to allow for easy comparison to the VAE later.

```

# we will sample N embeddings, then decode and visualize them
def vis_samples(model):
    ##### TODO
    #####
    # Sample embeddings from a diagonal unit Gaussian distribution and
    # decode them
    # using the model.
    #
    # HINT: The sampled embeddings should have shape [batch_size, nz].
    Diagonal unit
    # Gaussians have mean 0 and a covariance matrix with ones on
    the diagonal
    # and zeros everywhere else.
    #

```

```


# HINT: If you are unsure whether you sampled the correct
distribution, you can #
# sample a large batch and compute the empirical mean and
variance using the #
# .mean() and .var() functions.
#
# HINT: You can directly use model.decoder() to decode the samples.
#

#####
#####
sampled_embeddings = torch.randn(64, 64).to(device) # sample batch
of embedding from prior
decoded_samples = ae_model.decoder(sampled_embeddings) #
decoder output images for sampled embeddings
##### END TODO
#####

fig = plt.figure(figsize = (10, 10))
ax1 = plt.subplot(111)
ax1.imshow(torchvision.utils.make_grid(decoded_samples[:16], nrow=4,
pad_value=1.)\
            .data.cpu().numpy().transpose(1, 2, 0), cmap='gray')
plt.show()

vis_samples(ae_model)

```



**Inline Question: Describe your observations, why do you think they occur?**  
[2pt] \ (please limit your answer to <150 words) \ **Answer:**

The decoding model has learned the representation of the MNIST dataset in such a way that it can reconstruct the original image from this reduced dimensional space with minimal error. Here it is trying to reconstruct the random values to a digit, and we can see that it resembles a few digits, like 8 and 3. These are the representations it has learned.

### 3. Variational Auto-Encoder (VAE)

Variational auto-encoders use a very similar architecture to deterministic auto-encoders, but are inherently stochastic models, i.e. we perform a stochastic sampling operation during the forward pass, leading to different different outputs every time we run the network for the same input. This sampling is required to optimize the VAE objective also known as the evidence lower bound (ELBO):

$$p(x) > E_{z \sim q(z|x)} p(x \vee z) - D_{\text{KL}}(q(z|x), p(z))$$

Here,  $D_{\text{KL}}(q, p)$  denotes the Kullback-Leibler (KL) divergence between the posterior distribution  $q(z|x)$ , i.e. the output of our encoder, and  $p(z)$ , the prior over the embedding variable  $z$ , which we can choose freely.

For simplicity, we will again choose a unit Gaussian prior. The left term is the reconstruction term we already know from training the auto-encoder. When assuming a Gaussian output distribution for both encoder  $q(z|x)$  and decoder  $p(x|z)$  the objective reduces to:

$$L_{\text{VAE}} = \sum_{x \sim D} (x - \hat{x})^2 - \beta \cdot D_{\text{KL}}(N(\mu_q, \sigma_q), N(0, I))$$

Here,  $\hat{x}$  is the reconstruction output of the decoder. In comparison to the auto-encoder objective, the VAE adds a regularizing term between the output of the encoder and a chosen prior distribution, effectively forcing the encoder output to not stray too far from the prior during training. As a result the decoder gets trained with samples that look pretty similar to samples from the prior, which will hopefully allow us to generate better images when using the VAE as a generative model and actually feeding it samples from the prior (as we have done for the AE before).

The coefficient  $\beta$  is a scalar weighting factor that trades off between reconstruction and regularization objective. We will investigate the influence of this factor in our experiments below.

If you need a refresher on VAEs you can check out this tutorial paper:  
<https://arxiv.org/abs/1606.05908>

### Reparametrization Trick

The sampling procedure inside the VAE's forward pass for obtaining a sample  $z$  from the posterior distribution  $q(z|x)$ , when implemented naively, is non-differentiable. However, since  $q(z|x)$  is parametrized with a Gaussian function, there is a simple trick to obtain a differentiable sampling operator, known as the *reparametrization trick*.

Instead of directly sampling  $z \sim N(\mu_q, \sigma_q)$  we can "separate" the network's predictions and the random sampling by computing the sample as:

$$z = \mu_q + \sigma_q * \epsilon, \epsilon \sim N(0, I)$$

Note that in this equation, the sample  $z$  is computed as a deterministic function of the network's predictions  $\mu_q$  and  $\sigma_q$  and therefore allows to propagate gradients through the sampling procedure.

**Note:** While in the equations above the encoder network parametrizes the standard deviation  $\sigma_q$  of the Gaussian posterior distribution, in practice we usually parametrize the **logarithm of the standard deviation**  $\log \sigma_q$  for numerical stability. Before sampling  $z$  we will then exponentiate the network's output to obtain  $\sigma_q$ .

## Defining the VAE Model [7pt]

```
def kl_divergence(mu1, log_sigma1, mu2, log_sigma2):
    """Computes KL[p|q] between two Gaussians defined by [mu,
    log_sigma]."""
    return (log_sigma2 - log_sigma1) + (torch.exp(log_sigma1) ** 2 +
    (mu1 - mu2) ** 2) \
        / (2 * torch.exp(log_sigma2) ** 2) - 0.5

class VAE(nn.Module):
    def __init__(self, nz, beta=1.0):
        super().__init__()
        self.beta = beta          # factor trading off between two loss
        components
        ##### TODO
        #####
        # Instantiate Encoder and Decoder.
        #
        # HINT: Remember that the encoder is now parametrizing a Gaussian
        distribution's #
        #         mean and log_sigma, so the dimensionality of the output
        embedding needs to #
        #         double.
        #
        #####
        #####
        self.encoder = Encoder(2 * nz)
        self.decoder = Decoder(nz)
        ##### END TODO
        #####

    def forward(self, x):
        ##### TODO
        #####
        # Implement the forward pass of the VAE.
        #
        # HINT: Your code should implement the following steps:
        #
        #         1. encode input x, split encoding into mean and
        log_sigma of Gaussian #
        #         2. sample z from inferred posterior distribution using
        #
        #         reparametrization trick
        #
        #         3. decode the sampled z to obtain the reconstructed
        image #
        #####
```



```

#####
    # encode input into posterior distribution  $q(z | x)$ 
    q = self.encoder(x)          # output of encoder (concatenated mean
    and log_sigma)

    # sample latent variable z with reparametrization
    z = q[:, :nz] + torch.rand_like(torch.log(q[:,
nz:]))*torch.exp(q[:, nz:]))    # batch of sampled embeddings
    # compute reconstruction
    reconstruction = self.decoder(z) # decoder reconstruction from
    embedding
    ##### END TODO
#####

    return {'q': q,
            'rec': reconstruction}

def loss(self, x, outputs):
    ##### TODO
    #####
    # Implement the loss computation of the VAE.
    #
    # HINT: Your code should implement the following steps:
    #
    # 1. compute the image reconstruction loss, similar to AE
    we use MSE loss #
    # 2. compute the KL divergence loss between the inferred
    posterior #
    # distribution and a unit Gaussian prior; you can use
    the provided #
    # function above for computing the KL divergence
    between two Gaussians #
    # parametrized by mean and log_sigma
    #
    # HINT: Make sure to compute the KL divergence in the correct
    order since it is #
    # not symmetric, ie.  $KL(p, q) \neq KL(q, p)$ 
    #

    #####
    #####
    # compute reconstruction loss
    rec = outputs['rec']
    q = outputs['q']
    rec_loss = nn.functional.mse_loss(x, rec)
    # compute KL divergence loss
    # kl_loss = torch.mean(torch.sum(kl_divergence(q[:, :nz],
    torch.log(q[:, nz:]), torch.zeros(nz).to(device),
    torch.zeros(nz).to(device)).nan_to_num(), axis=1))# make sure that
    this is a scalar, not a vector / array

```

```

        kl_loss =
torch.mean(torch.sum(kl_divergence(torch.zeros(nz).to(device),
torch.zeros(nz).to(device), q[:, :nz], q[:, nz:]).nan_to_num(),
axis=1))# make sure that this is a scalar, not a vector / array
##### END TODO
#####

    # return weihgted objective
    return rec_loss + self.beta * kl_loss, \
        {'rec_loss': rec_loss, 'kl_loss': kl_loss}

def reconstruct(self, x):
    """Use mean of posterior estimate for visualization
    reconstruction."""
    ##### TODO
    #####
    # This function is used for visualizing reconstructions of our VAE
    model. To      #
    # obtain the maximum likelihood estimate we bypass the sampling
    procedure of the #
    # inferred latent and instead directly use the mean of the
    inferred posterior. #
    # HINT: encode the input image and then decode the mean of the
    posterior to obtain #
    #         the reconstruction.
    #

#####
#####
    q = self.encoder(x)
    reconstruction = self.decoder(q[:, :nz])
    ##### END TODO
    #####
    return reconstruction

```

## Setting up the VAE Training Loop [4pt]

Let's start training the VAE model! We will first verify our implementation by setting  $\beta=0$ .

```

learning_rate = 1e-3
nz = 64

```

```

##### TODO
#####
# Tune the beta parameter to obtain good VAE training results.
However, for the #
# initial experiments leave beta = 0 in order to verify our
implementation. #
#####
#####

```

```

epochs = 20          # using 5 epochs is sufficient for the first two
                      # experiments
                      # for the experiment where you tune beta, 20 epochs
                      # are appropriate
beta = 0.0001
##### END TODO
#####

# build VAE model
vae_model = VAE(nz, beta).to(device)    # transfer model to GPU if
available
vae_model = vae_model.train()    # set model in train mode (eg
batchnorm params get updated)

# build optimizer and loss function
##### TODO
#####
# Build the optimizer for the vae_model. We will again use the Adam
optimizer with #
# the given learning rate and otherwise default parameters.
#
#####
#####
opt = torch.optim.Adam(vae_model.parameters(), lr=learning_rate)
##### END TODO
#####

train_it = 0
rec_loss, kl_loss = [], []
for ep in range(epochs):
    print("Run Epoch {}".format(ep))
    ##### TODO
    #####
    # Implement the main training loop for the VAE model.
    #
    # HINT: Your training loop should sample batches from the data
    loader, run the #
    # forward pass of the VAE, compute the loss, perform the
    backward pass and #
    # perform one gradient step with the optimizer.
    #
    # HINT: Don't forget to erase old gradients before performing the
    backward pass. #
    # HINT: This time we will use the loss() function of our model for
    computing the #
    # training loss. It outputs the total training loss and a dict
    containing #
    # the breakdown of reconstruction and KL loss.
    #

```

```
#####
#####
for (x, _) in mnist_data_loader:
    x = x.to(device)
    forward_outs = vae_model.forward(x)
    total_loss, losses = vae_model.loss(x, forward_outs)
    total_loss = total_loss.cpu()
    vae_model.zero_grad()
    opt.zero_grad()
    total_loss.backward()
    losses['rec_loss'] = losses['rec_loss'].detach().cpu()
    losses['kl_loss'] = losses['kl_loss'].detach().cpu()
    opt.step()
    # add VAE training loop commands here
    # ...
    ##### END TODO
#####
```

```
    rec_loss.append(losses['rec_loss']);
kl_loss.append(losses['kl_loss'])
if train_it % 100 == 0:
    print("It {}: Total Loss: {}, \t Rec Loss: {},\t KL Loss: {}"\
          .format(train_it, total_loss, losses['rec_loss'],
                  losses['kl_loss']))
    train_it += 1
```

```
print("Done!")
```

```
# log the loss training curves
fig = plt.figure(figsize = (10, 5))
ax1 = plt.subplot(121)
ax1.plot(rec_loss)
ax1.title.set_text("Reconstruction Loss")
ax2 = plt.subplot(122)
ax2.plot(kl_loss)
ax2.title.set_text("KL Loss")
plt.show()
```

Run Epoch 0

```
It 0: Total Loss: 0.2337312400341034, Rec Loss: 0.23162946105003357,
      KL Loss: 21.01778221130371
```

```
It 100: Total Loss: 0.036748774349689484, Rec Loss:
0.03539176657795906, KL Loss: 13.570079803466797
```

```
It 200: Total Loss: 0.024348070845007896, Rec Loss:
0.02283547632396221, KL Loss: 15.125946998596191
```

```
It 300: Total Loss: 0.02180352248251438, Rec Loss:
0.01993614435195923, KL Loss: 18.673782348632812
```

```
It 400: Total Loss: 0.018534326925873756, Rec Loss:
0.016756225377321243, KL Loss: 17.781015396118164
```

```
It 500: Total Loss: 0.01699286326766014, Rec Loss:
```

0.015122673474252224, KL Loss: 18.701904296875  
It 600: Total Loss: 0.01705174706876278, Rec Loss:  
0.015139229595661163, KL Loss: 19.125171661376953  
It 700: Total Loss: 0.01530380453914404, Rec Loss:  
0.013450547121465206, KL Loss: 18.53257179260254  
It 800: Total Loss: 0.016046034172177315, Rec Loss:  
0.014114723540842533, KL Loss: 19.3131046295166  
It 900: Total Loss: 0.013427069410681725, Rec Loss:  
0.011467210948467255, KL Loss: 19.59857940673828

Run Epoch 1

It 1000: Total Loss: 0.015417631715536118, Rec Loss:  
0.013443171977996826, KL Loss: 19.74459457397461  
It 1100: Total Loss: 0.014303203672170639, Rec Loss:  
0.012274106964468956, KL Loss: 20.290969848632812  
It 1200: Total Loss: 0.013512368313968182, Rec Loss:  
0.011571221053600311, KL Loss: 19.411470413208008  
It 1300: Total Loss: 0.014181634411215782, Rec Loss:  
0.012233473360538483, KL Loss: 19.481609344482422  
It 1400: Total Loss: 0.01398109644651413, Rec Loss:  
0.011965928599238396, KL Loss: 20.151676177978516  
It 1500: Total Loss: 0.0135706327855587, Rec Loss:  
0.011567069217562675, KL Loss: 20.0356388092041  
It 1600: Total Loss: 0.01250480581074953, Rec Loss:  
0.010474496521055698, KL Loss: 20.303096771240234  
It 1700: Total Loss: 0.012361840344965458, Rec Loss:  
0.010396509431302547, KL Loss: 19.65330696105957  
It 1800: Total Loss: 0.013560134917497635, Rec Loss:  
0.011530271731317043, KL Loss: 20.29863739013672

Run Epoch 2

It 1900: Total Loss: 0.012573613785207272, Rec Loss:  
0.010611371137201786, KL Loss: 19.622426986694336  
It 2000: Total Loss: 0.012237310409545898, Rec Loss:  
0.01025467086583376, KL Loss: 19.826400756835938  
It 2100: Total Loss: 0.011951389722526073, Rec Loss:  
0.010073976591229439, KL Loss: 18.774131774902344  
It 2200: Total Loss: 0.013662638142704964, Rec Loss:  
0.011595742776989937, KL Loss: 20.668954849243164  
It 2300: Total Loss: 0.011914394795894623, Rec Loss:  
0.009933087974786758, KL Loss: 19.813068389892578  
It 2400: Total Loss: 0.01227217074483633, Rec Loss:  
0.010211063548922539, KL Loss: 20.611074447631836  
It 2500: Total Loss: 0.011185338720679283, Rec Loss:  
0.009228352457284927, KL Loss: 19.569866180419922  
It 2600: Total Loss: 0.013756802305579185, Rec Loss:  
0.011653892695903778, KL Loss: 21.029094696044922  
It 2700: Total Loss: 0.01224315632134676, Rec Loss:  
0.010243695229291916, KL Loss: 19.994609832763672  
It 2800: Total Loss: 0.011746814474463463, Rec Loss:  
0.009737318381667137, KL Loss: 20.0949649810791

Run Epoch 3

It 2900: Total Loss: 0.01197094563394785, Rec Loss: 0.009961074218153954, KL Loss: 20.098712921142578  
It 3000: Total Loss: 0.011041415855288506, Rec Loss: 0.009078983217477798, KL Loss: 19.62432861328125  
It 3100: Total Loss: 0.012746147811412811, Rec Loss: 0.010597671382129192, KL Loss: 21.484764099121094  
It 3200: Total Loss: 0.010808959603309631, Rec Loss: 0.008815997280180454, KL Loss: 19.929628372192383  
It 3300: Total Loss: 0.012135517783463001, Rec Loss: 0.01000713836401701, KL Loss: 21.283798217773438  
It 3400: Total Loss: 0.012651904486119747, Rec Loss: 0.01066006813198328, KL Loss: 19.91836166381836  
It 3500: Total Loss: 0.011761651374399662, Rec Loss: 0.00978136621415615, KL Loss: 19.802852630615234  
It 3600: Total Loss: 0.013106169179081917, Rec Loss: 0.011067613959312439, KL Loss: 20.385547637939453  
It 3700: Total Loss: 0.010718959383666515, Rec Loss: 0.008789443410933018, KL Loss: 19.295164108276367

Run Epoch 4

It 3800: Total Loss: 0.011049985885620117, Rec Loss: 0.008932968601584435, KL Loss: 21.170177459716797  
It 3900: Total Loss: 0.011663142591714859, Rec Loss: 0.009615457616746426, KL Loss: 20.476852416992188  
It 4000: Total Loss: 0.011229131370782852, Rec Loss: 0.009225407615303993, KL Loss: 20.037242889404297  
It 4100: Total Loss: 0.012598467990756035, Rec Loss: 0.01062050648033619, KL Loss: 19.779621124267578  
It 4200: Total Loss: 0.0114980423822999, Rec Loss: 0.00949908047914505, KL Loss: 19.98961639404297  
It 4300: Total Loss: 0.01127099059522152, Rec Loss: 0.009196130558848381, KL Loss: 20.748598098754883  
It 4400: Total Loss: 0.011609885841608047, Rec Loss: 0.009655353613197803, KL Loss: 19.54532814025879  
It 4500: Total Loss: 0.010655476711690426, Rec Loss: 0.00867842324078083, KL Loss: 19.770538330078125  
It 4600: Total Loss: 0.011383214965462685, Rec Loss: 0.009324452839791775, KL Loss: 20.587627410888672

Run Epoch 5

It 4700: Total Loss: 0.010912141762673855, Rec Loss: 0.008901283144950867, KL Loss: 20.10858726501465  
It 4800: Total Loss: 0.01162390224635601, Rec Loss: 0.00957073736935854, KL Loss: 20.531654357910156  
It 4900: Total Loss: 0.010170238092541695, Rec Loss: 0.008329479955136776, KL Loss: 18.40758514404297  
It 5000: Total Loss: 0.011235630139708519, Rec Loss: 0.009235206991434097, KL Loss: 20.00423240661621  
It 5100: Total Loss: 0.011403016746044159, Rec Loss: 0.009313545189797878, KL Loss: 20.894710540771484  
It 5200: Total Loss: 0.010562351904809475, Rec Loss: 0.008551323786377907, KL Loss: 20.110279083251953

It 5300: Total Loss: 0.011084072291851044, Rec Loss: 0.009131333790719509, KL Loss: 19.527389526367188  
It 5400: Total Loss: 0.010685810819268227, Rec Loss: 0.008706529624760151, KL Loss: 19.79280662536621  
It 5500: Total Loss: 0.011725720949470997, Rec Loss: 0.009676782414317131, KL Loss: 20.48938751220703  
It 5600: Total Loss: 0.010762940160930157, Rec Loss: 0.00870999600738287, KL Loss: 20.52943992614746

Run Epoch 6

It 5700: Total Loss: 0.010728634893894196, Rec Loss: 0.00882148090749979, KL Loss: 19.071537017822266  
It 5800: Total Loss: 0.011182880029082298, Rec Loss: 0.009180882014334202, KL Loss: 20.019981384277344  
It 5900: Total Loss: 0.011242968030273914, Rec Loss: 0.009294232353568077, KL Loss: 19.487354278564453  
It 6000: Total Loss: 0.011706531047821045, Rec Loss: 0.00965104065835476, KL Loss: 20.554901123046875  
It 6100: Total Loss: 0.010770891793072224, Rec Loss: 0.008772864006459713, KL Loss: 19.98027992248535  
It 6200: Total Loss: 0.010298667475581169, Rec Loss: 0.008425189182162285, KL Loss: 18.734783172607422  
It 6300: Total Loss: 0.010512295179069042, Rec Loss: 0.008560956455767155, KL Loss: 19.513385772705078  
It 6400: Total Loss: 0.010647623799741268, Rec Loss: 0.008569744415581226, KL Loss: 20.778797149658203  
It 6500: Total Loss: 0.009606734849512577, Rec Loss: 0.007599283009767532, KL Loss: 20.07451629638672

Run Epoch 7

It 6600: Total Loss: 0.009882181882858276, Rec Loss: 0.007866312749683857, KL Loss: 20.158687591552734  
It 6700: Total Loss: 0.011033924296498299, Rec Loss: 0.009002123028039932, KL Loss: 20.318016052246094  
It 6800: Total Loss: 0.011210756376385689, Rec Loss: 0.009148713201284409, KL Loss: 20.620433807373047  
It 6900: Total Loss: 0.010074814781546593, Rec Loss: 0.008128092624247074, KL Loss: 19.46722412109375  
It 7000: Total Loss: 0.009631154127418995, Rec Loss: 0.00773963239043951, KL Loss: 18.91521453857422  
It 7100: Total Loss: 0.00967472791671753, Rec Loss: 0.007734966930001974, KL Loss: 19.397615432739258  
It 7200: Total Loss: 0.01078040711581707, Rec Loss: 0.008755271323025227, KL Loss: 20.251361846923828  
It 7300: Total Loss: 0.010695084929466248, Rec Loss: 0.008751566521823406, KL Loss: 19.4351806640625  
It 7400: Total Loss: 0.010861076414585114, Rec Loss: 0.00889766775071621, KL Loss: 19.634092330932617  
It 7500: Total Loss: 0.010362161323428154, Rec Loss: 0.008475693874061108, KL Loss: 18.864673614501953

Run Epoch 8

It 7600: Total Loss: 0.011189166456460953, Rec Loss:

0.009079241193830967, KL Loss: 21.09925079345703  
It 7700: Total Loss: 0.010027027688920498, Rec Loss:  
0.008037402294576168, KL Loss: 19.89625358581543  
It 7800: Total Loss: 0.009904044680297375, Rec Loss:  
0.007940786890685558, KL Loss: 19.632579803466797  
It 7900: Total Loss: 0.009269644506275654, Rec Loss:  
0.00732524786144495, KL Loss: 19.443967819213867  
It 8000: Total Loss: 0.011116919107735157, Rec Loss:  
0.009120163507759571, KL Loss: 19.96755599975586  
It 8100: Total Loss: 0.009996791370213032, Rec Loss:  
0.007930874824523926, KL Loss: 20.659164428710938  
It 8200: Total Loss: 0.009633170440793037, Rec Loss:  
0.00773466844111681, KL Loss: 18.985021591186523  
It 8300: Total Loss: 0.009711318649351597, Rec Loss:  
0.007771733216941357, KL Loss: 19.395851135253906  
It 8400: Total Loss: 0.00980121549218893, Rec Loss:  
0.007797365076839924, KL Loss: 20.038501739501953

Run Epoch 9

It 8500: Total Loss: 0.010241026990115643, Rec Loss:  
0.008257930167019367, KL Loss: 19.830970764160156  
It 8600: Total Loss: 0.01074901595711708, Rec Loss:  
0.008725578896701336, KL Loss: 20.234376907348633  
It 8700: Total Loss: 0.011062851175665855, Rec Loss:  
0.009063187055289745, KL Loss: 19.996646881103516  
It 8800: Total Loss: 0.010644344612956047, Rec Loss:  
0.008598100394010544, KL Loss: 20.46244239807129  
It 8900: Total Loss: 0.009751505218446255, Rec Loss:  
0.0077975899912416935, KL Loss: 19.539155960083008  
It 9000: Total Loss: 0.010029444471001625, Rec Loss:  
0.008111625909805298, KL Loss: 19.178184509277344  
It 9100: Total Loss: 0.010985367931425571, Rec Loss:  
0.008956646546721458, KL Loss: 20.287212371826172  
It 9200: Total Loss: 0.009876340627670288, Rec Loss:  
0.007843519560992718, KL Loss: 20.328210830688477  
It 9300: Total Loss: 0.009982575662434101, Rec Loss:  
0.008055259473621845, KL Loss: 19.273160934448242

Run Epoch 10

It 9400: Total Loss: 0.010141980834305286, Rec Loss:  
0.008086148649454117, KL Loss: 20.558320999145508  
It 9500: Total Loss: 0.010368820279836655, Rec Loss:  
0.00838407315313816, KL Loss: 19.847469329833984  
It 9600: Total Loss: 0.010149510577321053, Rec Loss:  
0.00813536997884512, KL Loss: 20.141407012939453  
It 9700: Total Loss: 0.009718629531562328, Rec Loss:  
0.007757882121950388, KL Loss: 19.60747718811035  
It 9800: Total Loss: 0.008868706412613392, Rec Loss:  
0.007033166475594044, KL Loss: 18.355396270751953  
It 9900: Total Loss: 0.009782112203538418, Rec Loss:  
0.00785546749830246, KL Loss: 19.26644515991211  
It 10000: Total Loss: 0.009920748881995678, Rec Loss:



0.008012413047254086, KL Loss: 19.083362579345703  
It 10100: Total Loss: 0.011020917445421219, Rec Loss:  
0.008964565582573414, KL Loss: 20.563518524169922  
It 10200: Total Loss: 0.009082206524908543, Rec Loss:  
0.007148674223572016, KL Loss: 19.33531951904297  
It 10300: Total Loss: 0.010194560512900352, Rec Loss:  
0.008293364197015762, KL Loss: 19.011959075927734

Run Epoch 11

It 10400: Total Loss: 0.00940373633056879, Rec Loss:  
0.007542531006038189, KL Loss: 18.612051010131836  
It 10500: Total Loss: 0.009233261458575726, Rec Loss:  
0.007266642991453409, KL Loss: 19.666183471679688  
It 10600: Total Loss: 0.00848186295479536, Rec Loss:  
0.006556359119713306, KL Loss: 19.255041122436523  
It 10700: Total Loss: 0.009250662289559841, Rec Loss:  
0.007337313145399094, KL Loss: 19.13349151611328  
It 10800: Total Loss: 0.00948668085038662, Rec Loss:  
0.007473149802535772, KL Loss: 20.135305404663086  
It 10900: Total Loss: 0.010043583810329437, Rec Loss:  
0.008041223511099815, KL Loss: 20.023609161376953  
It 11000: Total Loss: 0.010067366994917393, Rec Loss:  
0.008146263659000397, KL Loss: 19.211030960083008  
It 11100: Total Loss: 0.010544667951762676, Rec Loss:  
0.00850841123610735, KL Loss: 20.362565994262695  
It 11200: Total Loss: 0.009751818142831326, Rec Loss:  
0.007756262086331844, KL Loss: 19.9555606842041

Run Epoch 12

It 11300: Total Loss: 0.009799700230360031, Rec Loss:  
0.007853757590055466, KL Loss: 19.459426879882812  
It 11400: Total Loss: 0.010336529463529587, Rec Loss:  
0.0083177974447608, KL Loss: 20.187318801879883  
It 11500: Total Loss: 0.00938414130359888, Rec Loss:  
0.00731789181008935, KL Loss: 20.662498474121094  
It 11600: Total Loss: 0.010069755837321281, Rec Loss:  
0.008138229139149189, KL Loss: 19.315269470214844  
It 11700: Total Loss: 0.009680209681391716, Rec Loss:  
0.007745346054434776, KL Loss: 19.34864044189453  
It 11800: Total Loss: 0.009413301944732666, Rec Loss:  
0.007457858417183161, KL Loss: 19.554431915283203  
It 11900: Total Loss: 0.009493906982243061, Rec Loss:  
0.007567024789750576, KL Loss: 19.26882553100586  
It 12000: Total Loss: 0.009318679571151733, Rec Loss:  
0.007405176293104887, KL Loss: 19.135028839111328  
It 12100: Total Loss: 0.00895090214908123, Rec Loss:  
0.007010401226580143, KL Loss: 19.405010223388672

Run Epoch 13

It 12200: Total Loss: 0.008735612034797668, Rec Loss:  
0.006883055437356234, KL Loss: 18.52556610107422  
It 12300: Total Loss: 0.009696024470031261, Rec Loss:  
0.007652694825083017, KL Loss: 20.433298110961914

It 12400: Total Loss: 0.009836899116635323, Rec Loss:  
0.007860262878239155, KL Loss: 19.76636505126953  
It 12500: Total Loss: 0.010244989767670631, Rec Loss:  
0.008260565809905529, KL Loss: 19.84423828125  
It 12600: Total Loss: 0.010686383582651615, Rec Loss:  
0.008695952594280243, KL Loss: 19.904312133789062  
It 12700: Total Loss: 0.010180667974054813, Rec Loss:  
0.008162700571119785, KL Loss: 20.17967414855957  
It 12800: Total Loss: 0.009238140657544136, Rec Loss:  
0.007156890816986561, KL Loss: 20.812503814697266  
It 12900: Total Loss: 0.009176095947623253, Rec Loss:  
0.007221668027341366, KL Loss: 19.544281005859375  
It 13000: Total Loss: 0.009474145248532295, Rec Loss:  
0.007469289004802704, KL Loss: 20.048564910888672  
It 13100: Total Loss: 0.009604974649846554, Rec Loss:  
0.00766359269618988, KL Loss: 19.413820266723633

Run Epoch 14

It 13200: Total Loss: 0.009036051109433174, Rec Loss:  
0.007228773087263107, KL Loss: 18.072784423828125  
It 13300: Total Loss: 0.009338753297924995, Rec Loss:  
0.007435858249664307, KL Loss: 19.028945922851562  
It 13400: Total Loss: 0.010226594284176826, Rec Loss:  
0.008227042853832245, KL Loss: 19.99551010131836  
It 13500: Total Loss: 0.009303375147283077, Rec Loss:  
0.007377778645604849, KL Loss: 19.255966186523438  
It 13600: Total Loss: 0.009831802919507027, Rec Loss:  
0.007921365089714527, KL Loss: 19.10437774658203  
It 13700: Total Loss: 0.010256019420921803, Rec Loss:  
0.00825662724673748, KL Loss: 19.993919372558594  
It 13800: Total Loss: 0.009492616169154644, Rec Loss:  
0.007542094215750694, KL Loss: 19.505220413208008  
It 13900: Total Loss: 0.009448716416954994, Rec Loss:  
0.007441073656082153, KL Loss: 20.076427459716797  
It 14000: Total Loss: 0.009207211434841156, Rec Loss:  
0.007303313352167606, KL Loss: 19.038978576660156

Run Epoch 15

It 14100: Total Loss: 0.009012767113745213, Rec Loss:  
0.007089658640325069, KL Loss: 19.2310848236084  
It 14200: Total Loss: 0.009808267466723919, Rec Loss:  
0.007831469178199768, KL Loss: 19.767982482910156  
It 14300: Total Loss: 0.009388512931764126, Rec Loss:  
0.007406635209918022, KL Loss: 19.818775177001953  
It 14400: Total Loss: 0.00828483421355486, Rec Loss:  
0.006444678641855717, KL Loss: 18.401559829711914  
It 14500: Total Loss: 0.008984452113509178, Rec Loss:  
0.007109673693776131, KL Loss: 18.747779846191406  
It 14600: Total Loss: 0.009112023748457432, Rec Loss:  
0.007162486203014851, KL Loss: 19.495376586914062  
It 14700: Total Loss: 0.009432138875126839, Rec Loss:  
0.0075379470363259315, KL Loss: 18.941917419433594

It 14800: Total Loss: 0.009679015725851059, Rec Loss:  
0.007672532461583614, KL Loss: 20.064836502075195  
It 14900: Total Loss: 0.008375903591513634, Rec Loss:  
0.006535575725138187, KL Loss: 18.403274536132812  
It 15000: Total Loss: 0.009418453089892864, Rec Loss:  
0.007482761051505804, KL Loss: 19.356922149658203

Run Epoch 16

It 15100: Total Loss: 0.00890457071363926, Rec Loss:  
0.0069479020312428474, KL Loss: 19.56668472290039  
It 15200: Total Loss: 0.009318983182311058, Rec Loss:  
0.007291019894182682, KL Loss: 20.279630661010742  
It 15300: Total Loss: 0.008803917095065117, Rec Loss:  
0.00691474974155426, KL Loss: 18.89167022705078  
It 15400: Total Loss: 0.008747617714107037, Rec Loss:  
0.006930071394890547, KL Loss: 18.175464630126953  
It 15500: Total Loss: 0.008583329617977142, Rec Loss:  
0.006689179688692093, KL Loss: 18.941497802734375  
It 15600: Total Loss: 0.008900785818696022, Rec Loss:  
0.0069357422180473804, KL Loss: 19.650434494018555  
It 15700: Total Loss: 0.009168310090899467, Rec Loss:  
0.007292360533028841, KL Loss: 18.759498596191406  
It 15800: Total Loss: 0.009695341810584068, Rec Loss:  
0.007654087617993355, KL Loss: 20.41254425048828  
It 15900: Total Loss: 0.009341772645711899, Rec Loss:  
0.007373183034360409, KL Loss: 19.685894012451172


Run Epoch 17

It 16000: Total Loss: 0.009631404653191566, Rec Loss:  
0.007655958645045757, KL Loss: 19.75446128845215  
It 16100: Total Loss: 0.00925498828291893, Rec Loss:  
0.007257900666445494, KL Loss: 19.970874786376953  
It 16200: Total Loss: 0.010051802732050419, Rec Loss:  
0.008001520298421383, KL Loss: 20.502826690673828  
It 16300: Total Loss: 0.009166950359940529, Rec Loss:  
0.007285760249942541, KL Loss: 18.81190299987793  
It 16400: Total Loss: 0.008989590220153332, Rec Loss:  
0.00714452937245369, KL Loss: 18.450605392456055  
It 16500: Total Loss: 0.008583547547459602, Rec Loss:  
0.006651151925325394, KL Loss: 19.32396125793457  
It 16600: Total Loss: 0.008448736742138863, Rec Loss:  
0.006481637712568045, KL Loss: 19.670995712280273  
It 16700: Total Loss: 0.009410754777491093, Rec Loss:  
0.007448025979101658, KL Loss: 19.627286911010742  
It 16800: Total Loss: 0.008248954080045223, Rec Loss:  
0.006393996998667717, KL Loss: 18.54957389831543

Run Epoch 18


It 16900: Total Loss: 0.00913525465875864, Rec Loss:  
0.0072027710266411304, KL Loss: 19.32483673095703  
It 17000: Total Loss: 0.009384341537952423, Rec Loss:  
0.007358510512858629, KL Loss: 20.258312225341797  
It 17100: Total Loss: 0.008171929977834225, Rec Loss:


0.006286827847361565, KL Loss: 18.85102081298828  
It 17200: Total Loss: 0.009632471948862076, Rec Loss:  
0.007718869484961033, KL Loss: 19.13602066040039  
It 17300: Total Loss: 0.008548183366656303, Rec Loss:  
0.006607479881495237, KL Loss: 19.407033920288086  
It 17400: Total Loss: 0.009624561294913292, Rec Loss:  
0.0076821427792310715, KL Loss: 19.42418670654297  
It 17500: Total Loss: 0.009428214281797409, Rec Loss:  
0.007521343417465687, KL Loss: 19.06870460510254  
It 17600: Total Loss: 0.009941276162862778, Rec Loss:  
0.007968501187860966, KL Loss: 19.727754592895508  
It 17700: Total Loss: 0.008030988276004791, Rec Loss:  
0.006184965837746859, KL Loss: 18.460229873657227  
It 17800: Total Loss: 0.008651576936244965, Rec Loss:  
0.006743073463439941, KL Loss: 19.085037231445312  
Run Epoch 19  
It 17900: Total Loss: 0.00953714083880186, Rec Loss:  
0.007608341984450817, KL Loss: 19.287992477416992  
It 18000: Total Loss: 0.009084980934858322, Rec Loss:  
0.007228220347315073, KL Loss: 18.567604064941406  
It 18100: Total Loss: 0.008538378402590752, Rec Loss:  
0.00660879397764802, KL Loss: 19.295848846435547  
It 18200: Total Loss: 0.008852812461555004, Rec Loss:  
0.006932985968887806, KL Loss: 19.19826889038086  
It 18300: Total Loss: 0.008395504206418991, Rec Loss:  
0.00651401374489069, KL Loss: 18.814905166625977  
It 18400: Total Loss: 0.009026801213622093, Rec Loss:  
0.007132979109883308, KL Loss: 18.93822479248047  
It 18500: Total Loss: 0.00881512276828289, Rec Loss:  
0.006911963690072298, KL Loss: 19.03158950805664  
It 18600: Total Loss: 0.008559603244066238, Rec Loss:  
0.006624417379498482, KL Loss: 19.351856231689453  
It 18700: Total Loss: 0.009245181456208229, Rec Loss:  
0.007343700621277094, KL Loss: 19.014808654785156  
Done!



Let's look at some reconstructions and decoded embedding samples!

```
# visualize VAE reconstructions and samples from the generative model  
vis_reconstruction(vae_model)  
vis_samples(vae_model)
```





**Inline Question: What can you observe when setting  $\beta=0$ ? Explain your observations! [3pt] \ (please limit your answer to <150 words) \ Answer:**

When  $B=0$ , we use only  $\log p(x|z)$  and ignore the prior divergence. We only use the reconstruction error. Hence we get the image shown above, without any disentanglement.

Let's repeat the same experiment for  $\beta=10$ , a very high value for the coefficient. You can modify the  $\beta$  value in the cell above and rerun it (it is okay to overwrite the outputs of the previous experiment, but **make sure to copy the visualizations of training curves, reconstructions and samples for  $\beta=0$  into your solution PDF before deleting them**).

**Inline Question: What can you observe when setting  $\beta=10$ ? Explain your observations! [3pt] \ (please limit your answer to <200 words) \ Answer:**

When  $B=10$ , we have a stronger constraint over the latent bottleneck. This greatly limits the representation capacity of  $z$  and results in further disentanglement, which in turn, results in the above image with very poor reconstruction. Optimal value exists between 0 and 10.

Now we can start tuning the beta value to achieve a good result. First describe what a "good result" would look like (focus what you would expect for reconstructions and sample quality).

**Inline Question: Characterize what properties you would expect for reconstructions (1pt) and samples (2pt) of a well-tuned VAE! [3pt]** \ (please limit your answer to <200 words) \ **Answer:**

A well tuned B-VAE will learn the disentangled representations while still having less reconstruction errors. They will however have lesser interpretability of the latent space.

### Tuning the $\beta$ -factor [5pt]

Now that you know what outcome we would like to obtain, try to tune  $\beta$  to achieve this result.

(logarithmic search in steps of 10x will be helpful, good results can be achieved after ~20 epochs of training). It is again okay to overwrite the results of the previous  $\beta=10$  experiment after copying them to the solution PDF.

**Your final notebook should include the visualizations of your best-tuned VAE.**

## 4. Embedding Space Interpolation [3pt]

As mentioned in the introduction, AEs and VAEs cannot only be used to generate images, but also to learn low-dimensional representations of their inputs. In this final section we will investigate the representations we learned with both models by **interpolating in embedding space** between different images. We will encode two images into their low-dimensional embedding representations, then interpolate these embeddings and reconstruct the result.

```
START_LABEL = 6
END_LABEL = 9
nz=64
```

```
def get_image_with_label(target_label):
    """Returns a random image from the training set with the requested
    digit."""
    for img_batch, label_batch in mnist_data_loader:
        for img, label in zip(img_batch, label_batch):
            if label == target_label:
                return img.to(device)

def interpolate_and_visualize(model, tag, start_img, end_img):
    """Encodes images and performs interpolation. Displays decodings."""
    model.eval()    # put model in eval mode to avoid updating batchnorm

    # encode both images into embeddings (use posterior mean for
```

```

interpolation)
z_start = model.encoder(start_img[None])[..., :nz]
z_end = model.encoder(end_img[None])[..., :nz]

# compute interpolated latents
N_INTER_STEPS = 5
z_inter = [z_start + i/N_INTER_STEPS * (z_end - z_start) for i in
range(N_INTER_STEPS)]


# decode interpolated embeddings (as a single batch)
img_inter = model.decoder(torch.cat(z_inter))

# reshape result and display interpolation
vis_imgs = torch.cat([start_img[None], img_inter, end_img[None]])
fig = plt.figure(figsize = (10, 10))
ax1 = plt.subplot(111)
ax1.imshow(torchvision.utils.make_grid(vis_imgs,
nrow=N_INTER_STEPS+2, pad_value=1.)\
.data.cpu().numpy().transpose(1, 2, 0), cmap='gray')
plt.title(tag)
plt.show()

# sample two training images with given labels
start_img = get_image_with_label(START_LABEL)
end_img = get_image_with_label(END_LABEL)

# visualize interpolations for AE and VAE models
interpolate_and_visualize(ae_model, "Auto-Encoder", start_img,
end_img)
interpolate_and_visualize(vae_model, "Variational Auto-Encoder",
start_img, end_img)

```





Repeat the experiment for different start / end labels and different samples. Describe your observations.

**Inline Question: Repeat the interpolation experiment with different start / end labels and multiple samples. Describe your observations! Focus on: \**

1. **How do AE and VAE embedding space interpolations differ? \**
2. **How do you expect these differences to affect the usefulness of the learned representation for downstream learning? \** (please limit your answer to <300 words)

**Answer:**

1. The latent space where the encoder encodes the input in an autoencoder may not be continuous and hence don't allow easy interpolation. On the other hand, the latent spaces of VAE are continuous and allow random sampling.
2. VAE learns both the mean and variance of the inputs in the latent space and hence can be used to generate new data from feeding random values to the decoder part. The autoencoder will try its best to generate an image which looks close to a digit, while the VAE can generate an image closer to a digit because it learns a distribution of possible latent inputs that can lead to the specific digit.

Hence, we can observe that the transition between the 2 classes is smoother in the case of VAE and very abrupt (you can clearly see 2 different digits in the case of 3&5) in the case of AE. This is due to the nature of the latent space.

## Submission PDF

As in assignment 1, please prepare a separate submission PDF for each problem. **Do not simply render your notebook as a pdf.** For this problem, please include the following plots & answers in a PDF called `problem_1_solution.pdf`:

1. Auto-encoder samples and AE sampling inline question answer.
2. VAE training curves, reconstructions and samples for:
  - $\beta=0$
  - $\beta=10$
  - your tuned  $\beta$  (also listing the tuned value for  $\beta$ )
1. Answers to all inline questions in VAE section (ie 4 inline questions).
2. Three representative interpolation comparisons that show AE and VAE embedding interpolation between the same images.
3. Answer to interpolation inline question.

Note that you still need to submit the jupyter notebook with all generated solutions. We will randomly pick submissions and check that the plots in the PDF and in the notebook are equivalent (except for those  $\beta=0/10$  plots that we allowed to overwrite).